

ATMA: LONG-CONTEXT LANGUAGE MODELING VIA POLAR ATTENTION AND GATED-DELTA COMPRESSION MEMORY

Anonymous authors

Paper under double-blind review

ABSTRACT

What determines whether a language model extrapolates beyond its training length? We study attention–memory interactions and checkpoint variability using **ATMA**, a 378M-parameter hybrid combining Polar Attention with gated-delta recurrent memory. Polar uses a normalized direction, a bounded participation-ratio magnitude, and learned length scaling. In a complete 120-cell, 1B-token factorial, memory improves Polar’s teacher-forced target-token retrieval at 64K in all 20 matched cells, by 47.8 percentage points on average. We then train matched NoPE, RoPE, and Polar models for 9.816B tokens at length 2K. The primary Polar checkpoint retains 62.7% exact FinePDFs retrieval at 8K and 34.4% target-token accuracy and 9.0% exact five-token accuracy overall at 256K, with all exact successes on synthetic haystacks. Two fresh paired runs retain a retrieval advantage over NoPE but reach only 0.0% and 3.83% exact accuracy. A post-hoc intervention caps the most retentive head in each checkpoint: primary NoPE and Polar 256K bits per byte fall from 8.097 to 1.595 and from 1.825 to 1.501. Both NoPE replications recover the likelihood improvement; the Polar replications have no comparable gate outlier and remain essentially unchanged by the cap. Short-context task accuracy and separately optimized Raven references expose additional quality and serving trade-offs. The evidence supports an attention–memory interaction and a checkpoint-specific retention failure, while leaving the incremental benefit of Polar’s full decomposition over temperature scaling unresolved.

1 INTRODUCTION

Large language models increasingly process code repositories, books, legal documents, and scientific papers that exceed their pretraining sequence length. Full softmax attention preserves token-level access in principle, but its probability mass can disperse as the key set grows (Nakanishi, 2025; Velickovic et al., 2025; Barbero et al., 2024). Sliding windows bound cost by removing older tokens from direct access. Recurrent sequence models instead keep fixed-size state but compress history. Long-context modeling therefore exposes a multi-objective trade-off among retrieval fidelity, sequence likelihood, reasoning, short-context quality, and serving cost.

Two questions organize this study: how does recurrent memory interact with the attention rule, and why can checkpoints with nearly identical short-context loss extrapolate differently? In the primary comparison, Polar delays the loss of exact natural-text associative recall relative to NoPE and RoPE, but every tested model eventually reaches zero. Fresh paired runs preserve a Polar retrieval advantage while showing substantially weaker extreme-length exact recall. We report the primary checkpoints and replications separately so that architectural trends and checkpoint-specific magnitudes remain distinguishable.

We present **ATMA**, a hybrid architecture that combines short gated convolutions with four **Polar Attention** layers as global mixers. Polar Attention represents “what matched” with a normalized direction and “how much matched” with a bounded participation-ratio channel. A learned null sink follows the extreme-value scale of irrelevant scores. Each global layer also contains gated-delta memory, providing a compressed recurrent history alongside full-context attention (Appendix A).




Figure 1: Primary untouched checkpoints by context length. Solid curves show the matched attention variants; dotted curves show separately optimized Raven references. **Left:** teacher-forced target-token retrieval. **Center:** BABILong exact match after adaptation. **Right:** mean fixed-target BPB (log scale; lower is better). Table 3 reports the fresh runs; the retention intervention is evaluated separately in Section 6.

We make three contributions. First, a complete 120-cell factorial measures how memory interacts with the attention core across regularizers, distractor alignment, and training windows. Second, matched 9.816B-token runs characterize retrieval, likelihood, reasoning, short-context quality, and serving; two fresh NoPE/Polar pairs test the retrieval and likelihood findings. Raven provides separately optimized recurrent-family references (Afzal et al., 2026). Third, we connect checkpoint-specific likelihood collapse to isolated near-unit retention gates through a selective inference-only intervention. The NoPE likelihood repair repeats in both fresh pairs, while Polar’s cap benefit does not. This localizes one extrapolation failure without establishing a general cap benefit or the training origin of the outlier.

2 RELATED WORK

Softmax attention (Vaswani et al., 2017) has quadratic compute and can lose selectivity as the number of keys grows. RoPE (Su et al., 2021) and ALiBi (Press et al., 2021) alter positional information while retaining dense simplex normalization. Scalable-Softmax, adaptive-temperature analyses, and sparse alternatives address dispersion or length-dependent sharpness (Nakanishi, 2025; Velickovic et al., 2025; Vasylenko et al., 2026). Threshold Differential Attention (TDA) uses a $\sqrt{\log n}$ threshold motivated by maxima of noise scores (Huang et al., 2026). Polar combines length-dependent temperature and extreme-value scaling with a differentiable null sink and a bounded magnitude channel. We compare Polar with matched NoPE and RoPE controls. Comparisons with the other methods are conceptual rather than empirical.

Linear attention, state-space models, and gated recurrences exchange full-resolution token access for compressed fixed-size history (Gu & Dao, 2023; Yang et al., 2024b;a; Behrouz et al., 2025). Near-unit gates create very long timescales and can enter saturation regimes (Gu et al., 2020). Stability-aware SSM parameterizations likewise distinguish stable eigenvalues from a robust distance to the stability boundary (Wang & Li, 2024). Stuffed Mamba directly studies an inability-to-forget failure and post-training interventions that restore forgetting (Chen et al., 2025). Our audit asks the corresponding checkpoint-level question for the gated-delta branch used here: a sigmoid keeps each retention value below one, but does not enforce a useful uniform margin from one. Raven combines routed sparse memory, decay dynamics, and gated-delta recurrence (Afzal et al., 2026). We treat it as a strong sub-quadratic baseline, but not as an attention ablation because it uses a different model family and optimizer. Hybrid models such as LFM2 combine local and global mixers for efficiency (Hasani et al., 2025; Allen-Zhu, 2025). ATMA follows this hybrid pattern while changing the global mixer and adding a recurrent memory stream.

3 METHOD

3.1 HYBRID ARCHITECTURE

ATMA is a 16-layer decoder whose block order repeats local–local–global–local four times, yielding 12 LFM2-style gated convolutional layers and global attention at blocks 3, 7, 11, and 15. All blocks use pre-normalization and squared-ReLU MLPs. The attention variants share hidden size 1024, eight query heads, two key-value heads, head dimension 128, tokenizer, data order, optimizer, schedule, and training budget. Only the attention core changes among NoPE, RoPE, and Polar. Canon depthwise causal convolutions are applied to q , k , and v before attention (Allen-Zhu, 2025).

3.2 POLAR ATTENTION

For query i and its $n_i = i + 1$ visible keys, let $\sigma_{ij} = q_i^\top k_j / \sqrt{d_k}$. Each head learns a length temperature and a virtual null logit,

$$t_i = 1 + \text{softplus}(\alpha) \log n_i, \quad \nu_i = b + \text{softplus}(\gamma) \sqrt{\log(n_i + 1)}. \quad (1)$$

The second term tracks the extreme-value scale of unrelated scores. Appending the null as a softmax competitor gives

$$w_i = \text{softmax}(t_i[\sigma_{i\bullet}, \nu_i]), \quad s_i = \sum_j w_{ij} v_j + w_{iN} v_{\text{null}}. \quad (2)$$

The *direction* channel is $c_i = s_i / \max(\|s_i\|_2, \epsilon)$, hence $\|c_i\|_2 \leq 1$. To retain bounded information about match multiplicity, real-key weights are renormalized as $\hat{w}_{ij} = w_{ij} / \sum_k w_{ik}$ and summarized by the participation ratio

$$n_{\text{eff},i} = \left(\sum_j \hat{w}_{ij}^2 \right)^{-1}, \quad m_{\text{eff},i} = n_{\text{eff},i} (1 - w_{iN}), \quad \mu_i = \tanh(\text{softplus}(\beta) \log(1 + m_{\text{eff},i})). \quad (3)$$

Thus $\mu_i \in [0, 1)$, while the direction channel removes radial scale. This is a channel decomposition, not an independence claim, because changing the match set can rotate c_i . The Polar output is a gated projection of c_i plus a learned projection of μ_i . Pre-integration tests rejected simpler choices. Additive soft counts leaked with length, fixed or relative floors were overtaken by noise tails, and an unbounded $\log m$ channel left the training range. The standalone prototype was not retained. Appendix C separates this design history from the reproducible integrated oracle and kernel tests. The promoted recipe disables the distractor-alignment loss tested in the factorial.

3.3 GATED-DELTA MEMORY

Each global layer also maintains a per-head recurrent state $M_t \in \mathbb{R}^{d_v \times d_k}$. With unit-normalized keys and queries, retention $\gamma_t = \sigma(W_\gamma x_t + b_\gamma)$, and write gate β_t , the decay-first, self-inclusive recurrence is

$$M_t = \gamma_t M_{t-1} (I - \beta_t k_t k_t^\top) + \beta_t v_t k_t^\top, \quad r_t = M_t q_t. \quad (4)$$

RMS-normalized r_t is gated, projected, and added to the Polar output. The projection is zero-initialized. Unit L_2 normalization is important. RMS-normalized keys have $\|k\|_2^2 = d_k$ and made the tested delta recurrence diverge, while unit keys keep the update eigenvalue along k within the gate-controlled range. Under bounded values and a uniformly contractive retention gate, Appendix A gives a sequence-length-independent state bound. The sigmoid guarantees $0 < \gamma_t < 1$ pointwise, but not the fixed margin $\gamma_t \leq \bar{\gamma} < 1$ assumed by that result. The implied half-life $H_{1/2} = \log(1/2) / \log \gamma$ diverges as $\gamma \rightarrow 1$. Section 6 audits this distinction empirically.

Figure 2 summarizes the two mechanisms and their controlled diagnostics. Appendix B defines every regularization mode in the factorial sweep, Appendix C gives the numerical oracle and streaming Triton implementation, and Appendix I details the kernel and serving optimizations.




Figure 2: Controlled ATMA diagnostics. **Left:** the effective match count maps to a bounded magnitude channel. **Right:** the tested gated-delta state stays flat with length, unlike an ungated Hebbian memory. These diagnostics illustrate the mechanism but do not prove network-level stability.

4 EXPERIMENTAL DESIGN

4.1 RECIPE SELECTION

Stage I trains all $3 \times 5 \times 2 \times 2 \times 2 = 120$ combinations of attention type (Polar, RoPE, NoPE), regularization, distractor loss, memory, and a 1024-token training window. Every model receives about 1B FineWeb-Edu tokens at length 2048 on an NVIDIA L4, and every evaluation uses full context. These are distinct factorial cells, not seed replicates. Stage I selects the promoted recipe and is not used for headline absolute comparisons.

4.2 MATCHED COMPARISON

Stage II trains 378.16–378.22M-parameter NoPE, RoPE, and Polar models for 9.816B tokens at length 2048 in the same L40S software and hardware environment. All three use the same Muon/AdamW parameter split. Raven Native (388.54M) and Atma-Raven-Titans (382.37M) use the same data, tokenizer, token budget, and device class, but use AdamW. In a 1B-token Raven Native pilot with the ATMA Muon schedule, validation loss improved to 3.76 nats at step 625, then became unstable and ended at 4.62. Retrieval remained at 0% through 64K. We therefore report the stable Raven models as a separate model-family/optimizer group. The single failed pilot does not establish incompatibility with every Muon configuration.

Evaluation. The primary retrieval grid has 24,000 model-trial evaluations across five models (4,800 per model): passkey/NIAH, synthetic/FinePDFs haystacks, eight lengths from 2K to 256K, three depths, and 50 trials per cell. Cues and digit targets are paired across lengths, depths, and models; all real-text trials reuse prefixes of one concatenated FinePDFs stream. We score five teacher-forced target tokens and exact five-token matches, rather than instruction-following generation. Fixed-target BPB scores one 256-token span in each of 18 documents (8 FinePDFs, 5 PG-19, 5 Proof-Pile), totaling 4,608 target tokens per model and length. Additional evaluations cover BABI-Long adaptation and greedy exact match (Kuratov et al., 2024), eight 2K tasks (22,939 examples per model), and single-sequence serving on one L40S. BABI-Long has ten held-out rows per task and length with custom splits and prompts. These are controlled diagnostics with limited independent text coverage; Appendix D gives aggregation and pairing details.

5 RESULTS

5.1 COMPONENT SELECTION

Table 1 summarizes matched contrasts from the full factorial rather than selected cells. At 64K, adding memory improves Polar in all 20 combinations of regularizer, distractor setting, and window

64K matched contrast	Mean Δ	Median Δ	Range	Positive
Polar: memory on – off	+47.8	+51.3	[+7.5, +92.5]	20/20
NoPE: memory on – off	+4.6	+1.9	[−15.0, +21.3]	11/20
RoPE: memory on – off	0.0	0.0	[0.0, 0.0]	0/20
Polar+memory – NoPE+memory	+41.7	+41.9	[0.0, +77.5]	19/20 (+1 tie)
Polar+memory – RoPE+memory	+50.3	+51.9	[+18.8, +92.5]	20/20

Table 1: Matched contrasts across the 120-cell, 1B-token factorial. Values are percentage-point changes in teacher-forced five-token accuracy at 64K. Each row holds all other factorial axes fixed. “Positive” excludes one numerical tie in the Polar–NoPE comparison.

Model	Retrieval 256K (%)		BABI 256K (%)	BPB	
	Token	Exact		2K	256K
NoPE	0.6	0.0	0	1.432	8.097
Polar	34.4	9.0	28	1.451	1.825
RoPE	0.0	0.0	14	1.439	2.512
Atma-Raven-Titans	10.0	0.0	38	1.525	1.551
Raven Native	20.3	0.0	46	1.581	1.597

Table 2: Primary untouched endpoints. The first three models form the matched attention group; the last two use Raven/AdamW. Retrieval averages tasks, suites, and depths; exact requires all five target tokens to be correct. BABILong is macro exact match after adaptation. BPB averages three fixed-target datasets (lower is better).

setting, by 47.8 points on average. Its effect on NoPE is small and mixed, and RoPE remains at zero. Among memory-enabled cells, Polar is never worse than NoPE and exceeds RoPE in all 20 comparisons. A training window reduces Polar+memory by 25.5 points on average and hurts 8 of 10 matched cells. Distractor alignment has a smaller, mixed effect, averaging −6.0 points and hurting 6 of 10 cells. We therefore promote full-context Polar plus memory without distractor alignment. The selected baseline cell reaches 92.5% token accuracy at 64K. Appendix E gives the complete matrix.

5.2 WHAT THE POLAR COMPONENTS ESTABLISH

A separate single-seed, 1B-token pilot tests Polar’s integrated components (Appendix C.2). At 64K, Full Polar reaches 77.50% target-token retrieval; removing the count channel retains 71.25%, fixing the null slope near zero improves it to 91.25%, and fixing the temperature gain near zero reduces it to 16.25%. Each length has 16 retrieval trials. These contrasts support learned length temperature and substantial retrieval through the direction-only path. They do not establish that the count channel or learned null slope is necessary. The full recipe remains the prespecified primary comparison, while its incremental benefit over temperature-scaled softmax and the null slope’s extrapolation calibration remain open.

5.3 A MULTI-OBJECTIVE FRONTIER

Table 2 reports the primary untouched checkpoints. At 2K, NoPE and Polar have nearly perfect target-token accuracy. At 256K, NoPE and RoPE fall to 0.6% and 0.0%, while Polar retains 34.4% averaged across tasks, suites, and depths. Exact five-token accuracy gives a stricter picture. On FinePDFs, Raven Native scores 0.0% throughout; Atma-Raven-Titans and RoPE reach zero at 4K and NoPE at 8K. Polar retains 93.0%, 67.3%, 62.7%, and 12.0% at 2K, 4K, 8K, and 16K, then reaches zero from 32K onward. Its synthetic exact accuracy is 65.7% at 64K and 18.0% at 256K. Thus its 9.0% overall exact result at 256K comes entirely from the synthetic suite. Appendix G gives the complete breakdown; Section 6 evaluates the cap separately.

Model	Training run	Retrieval 256K (%)		BPB 256K
		Token	Exact	
NoPE	Primary (mbs16)	0.60	0.00	8.097
NoPE	202701 (mbs4)	4.83	0.00	3.898
NoPE	202702 (mbs4)	12.13	0.00	5.256
Polar	Primary (mbs16)	34.37	9.00	1.825
Polar	202701 (mbs4)	16.97	0.00	1.644
Polar	202702 (mbs4)	20.90	3.83	1.760

Table 3: Untouched primary and fresh paired runs at 256K. All use 9.816B tokens and length 2K on L40S. Each fresh seed pairs NoPE and Polar; the primary checkpoints have no recorded common initialization seed and use a larger microbatch. These rows describe training-run variability, not an isolated seed effect. Exact successes at 256K are synthetic.




Figure 3: Untouched zero-shot accuracy on eight 2K controls (22,939 examples per model). The rightmost group gives unweighted task means. Raven variants are separately optimized references. Table 22 reports the inference-cap comparison.

5.4 REPLICATION OF RETRIEVAL AND LIKELIHOOD

Table 3 places the primary NoPE/Polar checkpoints beside two fresh seed-paired runs. Each pair shares initialization/data seeds and the same 9.816B-token budget; the fresh runs use microbatch size four versus sixteen for the primary comparison. At 256K, Polar exceeds paired NoPE token accuracy in both seeds, but its 16.97/20.90% token and 0.00/3.83% exact accuracy are substantially below the primary 34.37/9.00%. All exact successes are synthetic. Polar’s untouched BPB remains below paired NoPE in both seeds. Thus the direction of these two comparisons repeats, while the primary extreme-length retrieval magnitude does not. These are training-run contrasts, not estimates of a seed-only effect. Section 6 tests the cap in the same fresh pairs; Appendix F reports complete curves.

5.5 SHORT-CONTEXT QUALITY AND SERVING

Polar limits mean BPB degradation to $1.26\times$ at 256K, compared with $5.65\times$ for NoPE and $1.75\times$ for RoPE. This stability comes with a short-context cost. Mean accuracy over eight 2K controls is 43.29% for Polar, 45.15% for NoPE, and 44.60% for RoPE. Figure 3 shows that the gap varies by task. The 1.86-point mean gap to NoPE rules out a blanket quality claim. Appendix H gives the per-task values.

Table 4 shows the quality–systems trade-off. Raven Native reaches 46% BABILong at 256K and keeps BPB nearly flat. Its fixed recurrent state also avoids a sequence-length-dependent KV cache. At 128K, Atma-Raven-Titans decodes at 2.27 ms/token, compared with 16.3 ms/token for Polar. Polar instead has higher synthetic retrieval and training MFU. The BABILong samples are small. Simple Wilson intervals at 256K are 20.1–37.5% for Polar (28/100) and 36.6–55.7% for Raven Native (46/100). Because Raven also changes model family and optimizer, these rows show practical operating points rather than isolate one causal mechanism.

Group	Model	Base mean (%)	128K prefill (s)	Decode (ms/token)	Engine alloc. (GiB)	Train MFU (% , 2K)
Matched	NoPE	45.1	1.47	16.4	39.40	42.77
Matched	Polar	43.3	1.62	16.3	39.41	40.44
Matched	RoPE	44.6	1.47	16.5	39.41	42.45
Raven/AdamW	Atma-Raven-Titans	43.1	1.00	2.27	8.62	30.92
Raven/AdamW	Raven Native	43.0	1.36	3.27	6.46	21.53

Table 4: Short-context and systems trade-offs. Serving uses one sequence and one timing sample on one L40S. Engine allocation includes the attention backend’s preallocated KV pool (90% device-memory budget), so it is not the memory required by a single 128K sequence. Appendix I distinguishes pool allocation from occupied KV state.

6 CHECKPOINT VARIABILITY AND RETENTION-HORIZON AUDIT




Figure 4: NoPE 2K validation trajectories are nearly identical, yet the checkpoints’ 256K NLL values are 10.77, 9.90, and 4.07. The two L40S runs differ in microbatch size. The L4 run also differs in device class.

Figure 4 shows the trigger for a post-hoc audit. Three NoPE runs finish at 2.8126, 2.8128, and 2.8160 validation nats at 2K, yet reach 10.77, 9.90, and 4.07 NLL at 256K. The runs were not initialized from a recorded common seed, so the comparison establishes checkpoint variability, not a device-class treatment effect. Evaluation replays, a microbatch control, and deterministic fixtures ruled out several immediate explanations but did not identify the origin (Appendix L).

6.1 PARAMETER INSPECTION

We inspected the retention branch rather than whole-block gains. For each layer and query/memory head, we evaluate the learned zero-input operating point $\gamma_0 = \sigma(b_\gamma)$ and convert it to $H_{1/2} = \log(1/2)/\log \gamma_0$. Although the attention path has two KV heads, the memory gate has eight independently learned query/memory-head rows. The promoted NoPE checkpoint contains one block-2 head with a 21.01M-token half-life, and Polar contains one at 3.07M tokens. The corresponding L4 NoPE maximum is 185 tokens. Atma-Raven-Titans stays between 30.5 and 41.4 tokens. Because W_γ is nonzero, runtime retention still depends on activations. The scan is therefore a diagnostic operating point, not a complete trace. The outliers persist after BABILong fine-tuning, at 20.86M tokens for NoPE and 3.05M for Polar.

6.2 SINGLE-HEAD INTERVENTION AND FULL RE-EVALUATION

To test causality, we first conducted a paired multi-clamp sweep over candidate caps on the single highest zero-input operating point: checkpoint-local percentiles (p90, p99) and absolute half-life ceilings (h1:256, h1:512) across sequence lengths from 2K to 256K (Table 21). While p99 failed on saturated models because the 99th percentile was still dominated by the outlier ($H_{1/2} \approx$

484k tokens), the 256-token half-life cap

$$\gamma_{\max} = 2^{-1/256} = 0.997296 \quad (5)$$

gave the largest long-distance cross-entropy reduction while keeping 2K clean-loss drift within the 0.05-nat guardrail. In this same-process paired sweep, clean NLL at 256K fell from 12.003 to 1.138 for NoPE and from 1.252 to 1.057 for Polar. Needle cross-entropy fell from 12.20 to 6.05 and from 2.39 to 1.71, respectively. RoPE changed only from 6.16 to 6.02 and lost short-range retrieval, providing an empirical negative control.

We then used the selected 256-token ceiling to re-evaluate downstream tasks, retrieval, fixed-target BPB, and BABILong across the matched attention variants. Table 5 gives the intervention endpoints; Appendix J provides the complete before/after curves and numerical tables. All 15 jobs completed without an out-of-memory failure.

Model	Max $H_{1/2}$	Base mean	256K BPB	256K retrieval tok./exact	256K BABI
NoPE	21.01M	45.15 → 45.10	8.097 → 1.595	0.60/0.00 → 0.93/0.00	0 → 39
Polar	3.07M	43.29 → 43.25	1.825 → 1.501	34.37/9.00 → 47.07/16.33	28 → 42
RoPE	682	44.60 → 44.66	2.512 → 2.460	0.00/0.00 → 0.00/0.00	14 → 11

Table 5: Inference-only retention diagnostic. Results show archived baseline → one-head 256-token half-life cap. Base mean averages eight 2K tasks (%). Retrieval is token/exact accuracy (%), BABI is macro exact match, and BPB averages three fixed-target datasets. The cap is a probe, not a proposed trained model.

The intervention largely removes NoPE’s likelihood collapse and raises BABILong from 21/6/0/0/0% to 54/45/37/35/39% at 16/32/64/128/256K. However, NoPE’s exact natural-text retrieval remains zero beyond 8K and zero overall at 256K. State stability alone is therefore insufficient for exact natural-text recall. For Polar, the cap improves 256K token retrieval by 12.70 points, overall exact retrieval by 7.33 points, and BABILong by 14 points. Exact synthetic retrieval reaches 32.7%, while exact real-text retrieval reaches 91.3% at 8K and remains nonzero through 128K. Mean BPB degradation falls from $1.26\times$ to $1.03\times$. Mean accuracy on the eight 2K tasks changes by less than 0.06 points. RoPE does not recover retrieval, and its 256K BABILong score falls. The cap is therefore not a generic benchmark boost. For these promoted checkpoints, the intervention identifies excessive retention as a major mediator of NoPE and Polar degradation. The remaining retrieval gap shows that repairing this retention failure is insufficient to match Polar. It does not isolate which part of Polar’s attention rule contributes the advantage.

6.3 RETENTION INTERVENTION IN THE FRESH PAIRS

The fresh runs in Table 3 also received the same checkpoint-local cap. NoPE’s maximum zero-input half-lives were 135k and 5.78M tokens; the cap reduced 256K BPB from 3.898/5.256 to 1.506/1.490, but reduced token retrieval from 4.83/12.13% to 0.73/7.07%. Polar’s maxima were only 272 and 231 tokens, and its BPB and retrieval changed little. The NoPE likelihood repair therefore repeats in both pairs. Polar’s primary cap gain does not repeat, and neither architecture shows a general retrieval benefit from the cap. Appendix F gives all paired curves.

6.4 WHAT THIS AUDIT CANNOT ANSWER

The audit leaves several questions unresolved:

1. What training event produces the isolated outlier, and why did long horizons recur in both NoPE replications but not in the paired Polar checkpoints?
2. Would a bounded train-time parameterization prevent the outlier without removing useful long timescales? The inference cap does not answer this.
3. Why does NoPE exact retrieval remain zero on real text after likelihood and BABILong recover, and why does RoPE not benefit from the same cap?
4. Does Polar’s remaining natural-text retrieval gap at 256K improve with longer or mixed-domain training, or does it reflect distractor dispersion under unadapted extrapolation?

These questions delimit the current evidence rather than extend its claims. Resolving them requires interventions that separately vary architecture, optimizer, data, and execution environment.

7 LIMITATIONS AND CONCLUSION

The evidence covers one model scale and 2K pretraining, with two fresh NoPE/Polar pairs and no seed-matched RoPE replication. The factorial varies recipe settings rather than training seeds. Shared retrieval haystacks and 18 fixed BPB targets limit independent natural-text coverage; teacher-forced retrieval, adapted BABILong, and single-sample serving answer different questions. Exact FinePDFs retrieval reaches zero at extreme lengths in every tested model. Raven changes model family and optimizer. We have no matched 9.816B-token comparison with TDA or a softmax control sharing Polar’s learned temperature, and the component pilot does not establish an incremental benefit for the full decomposition. Its favorable fixed-null result requires replication and calibration measurements.

The recurrent bound assumes a uniform retention margin that the learned sigmoid does not enforce. The cap is a checkpoint intervention: it does not identify the outlier’s training origin or validate a bounded train-time formulation. Broad re-evaluation compares capped runs with pinned archived baselines; the smaller sweeps and fresh replications provide controlled pairing. We therefore distinguish replicated NoPE likelihood repair from checkpoint-specific Polar recovery.

ATMA establishes two empirical findings within this scope. Recurrent memory improves Polar retrieval across every matched factorial cell, and Polar retains a target-token advantage over NoPE in both fresh pairs, with weaker exact retrieval than in the primary checkpoint. Separately, a single retention-gate intervention repairs NoPE’s long-context likelihood collapse in the primary and both fresh runs despite similar short-context convergence. The component pilot and Raven references delimit the attribution and quality–systems trade-offs. Released configurations, logs, checkpoint manifests, benchmark matrices, clamp specifications, and proof skeletons make these findings auditable.

AI USE STATEMENT

Generative AI tools were used in this revision to assist diagnostic-script implementation, edit and structure the manuscript, and aggregate archived and newly produced evaluation logs. The authors ran the benchmark jobs on the listed checkpoints and hardware. AI was not used to generate evaluation datasets or train the reported models. The authors manually compare all AI-assisted mathematical text with the original derivations and proof artifacts, verify empirical values against source logs, and take responsibility for every claim and artifact in the submission. This statement describes the present version as required by the ICLR 2027 policy.

REPRODUCIBILITY STATEMENT

Model definitions and training settings are described in Sections 3–4. Appendix D specifies evaluation protocols and denominators, Appendix F reports the two seed-paired BPB/retrieval replications, Appendix J records the original parameter scan and full intervention, and Appendices E–K provide detailed tables and diagnostics. The anonymous artifact at <https://anonymous.4open.science/r/atma> contains configurations, logs, checkpoint manifests, benchmark payloads, clamp specifications, and machine-checked proof skeletons.

REFERENCES

- Arshia Afzal, Aviv Bick, Eric P. Xing, Volkan Cevher, and Albert Gu. Raven: High-recall sequence modeling with sparse memory routing. *arXiv preprint arXiv:2607.25357*, 2026. URL <https://arxiv.org/abs/2607.25357>.
- Habibullah Akbar. Weak-sigreg: Covariance regularization for stable deep learning. *arXiv preprint arXiv:2603.05924*, 2026.
- Zeyuan Allen-Zhu. Physics of language models: Part 4.1, architecture design and the magic of canon layers. *arXiv preprint arXiv:2512.17351*, 2025.
- Randall Balestriero and Yann LeCun. Lejepa: Provable and scalable self-supervised learning without the heuristics. *arXiv preprint arXiv:2511.08544*, 2025.
- Federico Barbero, Andrea Banino, Steven Kapturowski, Dharshan Kumaran, Joao G. M. Araujo, Alex Vitvitskyi, Razvan Pascanu, and Petar Velickovic. Transformers need glasses! information over-squashing in language tasks. In *Advances in Neural Information Processing Systems*, 2024.
- Ali Behrouz, Peilin Zhong, and Vahab Mirrokni. Titans: Learning to memorize at test time. *arXiv preprint arXiv:2501.00663*, 2025.
- Yingfa Chen, Xinrong Zhang, Shengding Hu, Xu Han, Zhiyuan Liu, and Maosong Sun. Stuffed mamba: Oversized states lead to the inability to forget. In *Second Conference on Language Modeling*, 2025.
- Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023.
- Albert Gu, Caglar Gulcehre, Thomas Paine, Matt Hoffman, and Razvan Pascanu. Improving the gating mechanism of recurrent neural networks. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pp. 3800–3809. PMLR, 2020.
- Ramin Hasani et al. Lfm2 technical report. *arXiv preprint arXiv:2511.23404*, 2025.
- Xingyue Huang, Xueying Ding, Mingxuan Ju, Yozen Liu, Neil Shah, and Tong Zhao. Threshold differential attention for sink-free, ultra-sparse, and non-dispersive language modeling. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*, 2026. URL <https://arxiv.org/abs/2601.12145>.

- Yuri Kuratov, Aydar Bulatov, Petr Anokhin, Ivan Rodkin, Dmitry Sorokin, Artyom Sorokin, and Mikhail Burtsev. Babilong: Testing the limits of llms with long context reasoning-in-a-haystack. In *Advances in Neural Information Processing Systems Datasets and Benchmarks Track*, 2024. URL <https://arxiv.org/abs/2406.10149>.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, pp. 611–626, 2023. doi: 10.1145/3600006.3613165.
- Ken M. Nakanishi. Scalable-softmax is superior for attention. *arXiv preprint arXiv:2501.19399*, 2025.
- Ofir Press, Noah A Smith, and Mike Lewis. Train short, test long: Attention with linear biases enables input length extrapolation. *arXiv preprint arXiv:2108.12409*, 2021.
- Jianlin Su et al. Roformer: Enhanced transformer with rotary position embedding. *arXiv preprint arXiv:2104.09864*, 2021.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, pp. 5998–6008. Curran Associates, Inc., 2017.
- Pavlo Vasylenko, Hugo Pitorro, Andre F. T. Martins, and Marcos Treviso. Long-context generalization with sparse attention. In *International Conference on Learning Representations*, 2026.
- Petar Velickovic, Christos Perivolaropoulos, Federico Barbero, and Razvan Pascanu. Softmax is not enough (for sharp size generalisation). In *International Conference on Machine Learning*, 2025.
- Shida Wang and Qianxiao Li. Stableness: Alleviating the curse of memory in state-space models through stable reparameterization. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 50766–50793. PMLR, 2024.
- Songlin Yang, Jan Kautz, and Ali Hatamizadeh. Gated delta networks: Improving mamba2 with delta rule. *arXiv preprint arXiv:2412.06464*, 2024a.
- Songlin Yang, Bailin Wang, Yu Zhang, Yikang Shen, and Yoon Kim. Parallelizing linear transformers with the delta rule over sequence length. *arXiv preprint arXiv:2406.06484*, 2024b.

A THEORETICAL FRAMING

These are sufficient-condition results. They establish noise-odds and norm bounds under explicit assumptions. They do not establish distributional invariance or replace the empirical evidence.

A.1 EXTREME-VALUE NULL FLOOR

Assumption 1 (Sub-Gaussian noise scores). *For irrelevant keys $j \in \mathcal{N}_n$, $|\mathcal{N}_n| \leq n$, scores X_j are mean-zero and v^2 -sub-Gaussian: $\mathbb{E} \exp(tX_j) \leq \exp(v^2 t^2/2)$ for all $t \in \mathbb{R}$.*

Let the Polar null floor be $\nu_n = b + s\sqrt{\log(n+1)}$ with $b, s \geq 0$.

Proposition 1 (Noise exceedances). *Under Assumption 1,*

$$\mathbb{E} \left[\sum_{j \in \mathcal{N}_n} \mathbf{1}\{X_j > \nu_n\} \right] \leq n \exp\left(-\frac{\nu_n^2}{2v^2}\right). \quad (6)$$

If $s \geq \sqrt{2}v$ this is $O(1)$; if the inequality is strict it decays polynomially.

Proof. Apply the sub-Gaussian tail bound to each key and sum. Since $b \geq 0$, $\nu_n^2 \geq s^2 \log(n+1)$. $\square$

Unlike TDA, Polar retains sub-threshold keys. With temperature $\theta_n = 1 + a \log n$, define the unnormalized noise-to-null odds $R_n = \sum_{j \in \mathcal{N}_n} \exp(\theta_n(X_j - \nu_n))$.

Proposition 2 (Vanishing null odds). *Suppose $s > c > \sqrt{2}v$ and the null floor is $s\sqrt{\log(n+1)}$ up to a nonnegative offset. With probability tending to one,*

$$R_n \leq n \exp\left(-(1 + a \log n)(s - c)\sqrt{\log n}\right). \quad (7)$$

For $a > 0$, the bound converges to zero faster than any inverse polynomial.

Proof. A union bound gives $\Pr[\max_j X_j > c\sqrt{\log(n+1)}] \leq n(n+1)^{-c^2/(2v^2)} \rightarrow 0$. On the complementary event, every noise score is at least $(s - c)\sqrt{\log n}$ below the floor. Bound each summand and sum over at most n keys. $\square$

A.2 BOUNDED CHANNELS AND RECURRENT STATE

Proposition 3 (Bounded Polar readout). *For $c_i = s_i / \max(\|s_i\|_2, \varepsilon)$, $\|c_i\|_2 \leq 1$. For $\mu_i = \tanh(\beta \log(1 + m_{\text{eff},i}))$, $\mu_i \in [0, 1)$ when $\beta, m_{\text{eff},i} \geq 0$. A distribution uniform over m real keys has participation ratio m .*

Proof. The first two claims follow from the denominator and range of $\tanh$. Uniform weights satisfy $(\sum_{j=1}^m (1/m)^2)^{-1} = m$. $\square$

Assumption 2 (Contractive memory). *For every t , $\|q_t\|_2 = \|k_t\|_2 = 1$, $\|v_t\|_2 \leq V$, $0 \leq \beta_t \leq 1$, and $0 \leq \gamma_t \leq \bar{\gamma} < 1$.*

Proposition 4 (Length-bounded state). *For recurrence equation 4,*

$$\|M_t\|_F \leq \bar{\gamma}^t \|M_0\|_F + \frac{V}{1 - \bar{\gamma}}. \quad (8)$$

Proof. $I - \beta_t k_t k_t^\top$ has spectral norm at most one and $\|\beta_t v_t k_t^\top\|_F \leq V$. Thus $\|M_t\|_F \leq \bar{\gamma} \|M_{t-1}\|_F + V$; unrolling gives the result. $\square$

The Lean artifact checks the finite-sum, exponential-monotonicity, bounded-channel composition, and invariant-set induction steps. It takes the sub-Gaussian tail, high-probability margin, and one-step contraction as premises. It is therefore a machine-checked proof skeleton, not end-to-end formal verification. The learned Polar conditions $b \geq 0$ and $s > c > \sqrt{2}v$ require an empirical head-wise audit before they can be asserted for a trained checkpoint. Likewise, $\gamma_t = \sigma(W_\gamma x_t + b_\gamma)$ establishes only $0 < \gamma_t < 1$ at any finite logit; it does not supply one sequence-uniform $\bar{\gamma} < 1$. Appendix J shows that the distinction is practically material: individual zero-input operating points can imply million-token half-lives, although that parameter-only scan does not by itself bound activation-conditioned runtime gates.

B REGULARIZATION MODES

The factorial grid varies a representation regularizer applied to the model’s signature stream. For every non-baseline mode, the sweep uses the same weight $\alpha_{\text{sig}} = 0.01$ and optimizes

$$\mathcal{L} = (1 - \alpha_{\text{sig}})\mathcal{L}_{\text{LM}} + \alpha_{\text{sig}}\mathcal{L}_{\text{reg}} + \lambda_{\text{dist}}\mathcal{L}_{\text{dist}}. \quad (9)$$

The weak and strong modes follow the two SIGReg families introduced by Weak-SIGReg (Akbar, 2026) and LeJEPa’s strong SIGReg objective (Balestrierio & LeCun, 2025). The discrete and zipfian modes are local variants that combine those ideas with stronger geometric priors. Table 6 defines all five sweep settings.

Mode	Meaning in the sweep
baseline	No signature-stream regularization; equivalently, $\alpha_{\text{sig}} = 0$.
weak	Weak SIGReg: center a random sketch of the token representations and match its covariance to the identity, $\ \text{Cov}(Sx) - I\ _F$. This encourages decorrelated, unit-variance features at low cost.
strong	Strong SIGReg: project representations onto random one-dimensional directions and match their empirical characteristic functions to the standard Gaussian characteristic function over a fixed frequency grid.
discrete	Local variant: normalize each representation to the sphere of radius $\sqrt{d}$, then apply weak covariance matching. This preserves whitening pressure while encouraging fixed-norm, discrete-code-like directions.
zipfian	Local variant: normalize directions for an orthogonality penalty, then match sorted representation magnitudes to a Zipf profile r^{-1} . This biases the signature stream toward heavy-tailed feature use rather than uniform energy allocation.

Table 6: Regularization modes in the 120-cell factorial sweep. Weak and strong are SIGReg-style covariance and distribution-matching objectives; discrete and zipfian are exploratory local variants.

C POLAR ATTENTION IMPLEMENTATION

C.1 DESIGN PROVENANCE AND SCOPE

Before integration, the Polar choices were developed by falsifying simpler alternatives in a standalone synthetic sandbox. Additive soft counts accumulated small per-key leakage with sequence length; thresholds based on a fixed number of standard deviations retained a constant fraction of noise keys; a fixed null logit was eventually overtaken by the maximum noise score; and an unbounded $\log m$ magnitude exposed the output projection to values beyond its training range. These failures motivated participation ratio, the $\sqrt{\log n}$ null floor, and the bounded $\tanh(\beta \log(1 + m))$ map. The standalone prototype was subsequently removed and is not part of the released artifact. We therefore use these tests as design provenance, not as model-scale component ablations or independently reproducible benchmark evidence. The integrated equations, materialized oracle, streaming implementation, and parity tests described below are retained.

C.2 INTEGRATED COMPONENT PILOT

We subsequently ran a matched network-level component study using the promoted Polar+memory recipe. Each 378.22M-parameter model received 1B tokens at length 2K with the same initialization and data seed; evaluation used the same documents, retrieval prompts, and seed. “Fixed null slope” sets only the nonnegative $\sqrt{\log(n+1)}$ coefficient in Equation 1 near zero; the null base and null vector remain learned. Table 7 reports the prespecified metrics. The clean and junk columns are NLL in nats/token on coherent FinePDFs and the concatenated FineWeb-Edu validation stream, respectively.

Polar variant	Val. NLL	Clean 2K	Clean 64K	Junk 2K	Junk 64K	Needle 2K / 64K (%)
Full Polar	3.1688	2.810	2.047	3.110	3.129	91.25 / 77.50
Direction only	3.1629	2.709	1.982	3.103	3.124	91.25 / 71.25
Constant magnitude	3.1717	2.734	1.924	3.113	3.134	92.50 / 16.25
Fixed null slope	3.1741	2.725	1.938	3.117	3.142	96.25 / 91.25
Fixed temperature gain	3.1695	2.707	1.967	3.111	3.155	93.75 / 16.25

Table 7: Single-seed, 1B-token Polar component attribution. Needle values are teacher-forced five-token accuracy over 16 trials per length. All variants share the same seed, but this design does not estimate variation across training seeds.

Ordinary language-model losses are tightly grouped: validation NLL spans only 0.0112 nats, and every intervention improves clean NLL relative to Full Polar at matched lengths. Retrieval is more diagnostic. Direction only retains most of the 64K accuracy (71.25% versus 77.50%), whereas constant magnitude and fixed temperature gain each fall to 16.25%. Thus the direction-only path carries

most of the retrieval signal, and the learned length temperature is important. An uninformative constant count signal is harmful, but this study does not establish that the count channel is always necessary.

Unexpectedly, fixing the length-dependent null slope near zero gives the best retrieval at every evaluated length and reaches 91.25% at 64K, while slightly worsening validation and junk-stream NLL. One plausible explanation is that a slope learned only through 2K over-rejects distant relevant keys when extrapolated, especially because this recipe disables the explicit distractor-alignment objective. The matched seed makes an independently luckier initialization an unlikely explanation for this contrast, but one training seed and 16 retrieval trials per length cannot establish a population-level effect. We therefore retain the prespecified Full Polar model for the confirmatory comparison and treat null calibration as unresolved. Future work should replicate the contrast across seeds and larger retrieval samples, measure null mass by head versus length, and test bounded or score-adaptive null formulations. The synthetic provenance above and this integrated pilot answer different questions: the former motivates the asymptotic design, while the latter shows that its learned finite-training calibration is not yet empirically optimal.

C.3 NUMERICAL ORACLE

The materialized PyTorch path forms the masked real-key logits and virtual null logit in FP32, applies a single softmax, and computes the direction and magnitude channels from the resulting weights. It is used as a numerical oracle for forward and backward tests, not as the production long-context path, because materializing the $T \times T$ score matrix is memory-prohibitive at the evaluated lengths.

C.4 STREAMING SUFFICIENT STATISTICS

The Triton kernel follows the online-max strategy used by FlashAttention-style kernels. For one query row, it streams K, V blocks and retains only four FP32 statistics: the running maximum M , real-key exponential sum L , squared exponential sum Q_2 , and unnormalized value sum S . If a new block changes the row maximum, $\alpha = \exp(M_{\text{old}} - M_{\text{new}})$ rescales each statistic as shown in Table 8.

State	Online update	Use after streaming
M	$M \leftarrow \max(M, \max a)$	Stable common logit origin
L	$L \leftarrow \alpha L + \sum_j p_j$	Real-key softmax mass
Q_2	$Q_2 \leftarrow \alpha^2 Q_2 + \sum_j p_j^2$	$n_{\text{eff}} = L^2 / Q_2$
S	$S \leftarrow \alpha S + \sum_j p_j v_j$	Direction numerator

Table 8: Streaming state for one Polar query row, where a is the current logit block and $p_j = \exp(a_j - M_{\text{new}})$. The squared accumulator must rescale by α^2 .

After the real keys are consumed, the virtual null is folded in as one additional softmax entry. Let $p_0 = \exp(t_i v_i - M)$ and $Z = L + p_0$. The kernel obtains

$$c_i = \text{normalize}(S + p_0 v_{\text{null}}), \quad n_{\text{eff},i} = \frac{L^2}{Q_2}, \quad m_{\text{eff},i} = n_{\text{eff},i} \frac{L}{Z}. \quad (10)$$

The common factor $1/Z$ cancels under radial normalization of the direction numerator. The production kernel evaluates the magnitude with $2\sigma(2x) - 1 = \tanh x$ and saves (M, L, Q_2, S) for the backward pass. Listing 1 gives the algorithmic core; pointer arithmetic, launch metadata, and backward code are omitted.

Listing 1: Schematic core of the block-streamed Triton forward pass.

```
@triton.jit
def polar_fwd_core(q, K, V, n_i, temp, nullv, beta,
                  BLOCK_N: tl.constexpr, DK: tl.constexpr):
    M = tl.full([BLOCK_M], -1e38, tl.float32)
    L = tl.zeros([BLOCK_M], tl.float32)
    Q2 = tl.zeros([BLOCK_M], tl.float32)
```



```

756 S = tl.zeros([BLOCK_M, DK], tl.float32)
757
758 for start in range(0, n_i, BLOCK_N):
759     k, v, valid = load_kv_block(K, V, start, n_i)
760     a = tl.where(valid, temp[:, None] * dot(q, k), -1e38)
761     M_new = tl.maximum(M, tl.max(a, axis=1))
762     alpha = tl.exp(M - M_new)
763     p = tl.where(valid, tl.exp(a - M_new[:, None]), 0.0)
764     L = alpha * L + tl.sum(p, axis=1)
765     Q2 = alpha * alpha * Q2 + tl.sum(p * p, axis=1)
766     S = alpha[:, None] * S + tl.dot(p, v)
767     M = M_new
768
769 # Fold in the learned virtual null as one extra entry.
770 M_new = tl.maximum(M, temp * nullv)
771 alpha = tl.exp(M - M_new)
772 L, Q2, S = alpha * L, alpha * alpha * Q2, alpha[:, None] * S
773 p_null = tl.exp(temp * nullv - M_new)
774 Z = L + p_null
775
776 direction = normalize(S + p_null[:, None] * v_null)
777 n_eff = L * L / tl.maximum(Q2, eps)
778 m_eff = n_eff * L / tl.maximum(Z, eps)
779 magnitude = 2 * tl.sigmoid(2 * beta * tl.log(1 + m_eff)) - 1
780 save_for_backward(M_new, L, Q2, S)

```

This removes the quadratic *score-matrix storage*. Full-context prefill still performs quadratic query-key arithmetic. The kernel is cross-checked against the materialized PyTorch oracle in FP32 accumulation before lower-precision outputs are accepted.

D EVALUATION PROTOCOLS

Stage I contains 120 approximately 370–378M-parameter runs trained for 1,900 steps of 524,288 tokens at sequence length 2,048 on FineWeb-Edu. Stage II contains three matched attention models trained for 18,722 such steps (9.816B tokens). The five promoted checkpoints and ten open-weight controls are evaluated with checkpoint-exact forward paths. Stage I cells vary factorial settings and are not independent seed replicates.

D.1 RAVEN OPTIMIZER CONTROL

The Raven rows use AdamW optimizer because the attempted matched-schedule control did not produce a usable baseline. In the 1B-token Raven Native pilot with the Atma Muon schedule, validation loss decreased from 10.83 to 3.76 nats by step 625, then rose to 6.88 at step 1,125 and ended at 4.62; teacher-forced retrieval was 0% at every tested length from 2K through 64K. The reported Raven Native and Atma-Raven-Titans checkpoints therefore use AdamW and remain a separate model-family/optimizer group. This is evidence about the tested Muon schedule, not a general claim that Raven cannot converge under any Muon tuning.

Retrieval crosses passkey and NIAH tasks, synthetic and FinePDFs haystacks, eight lengths from 2K through 256K, three depths, and 50 paired trials per cell. The scorer appends the complete five-token target and evaluates each position conditioned on the preceding ground-truth target tokens. Token accuracy is the fraction of correct per-position argmax predictions; exact accuracy requires all five positions to be correct. Both are teacher-forced diagnostics, not autoregressive generation. An older free-generation protocol was retired because all base checkpoints scored uniformly zero. BABILong uses answer-only fine-tuning on tasks qa1–qa10 at 0K/1K/2K and greedy exact match on ten held-out rows per task and length. Base-task controls use LAMBADA, HellaSwag, PIQA, WinoGrande, ARC-Easy, ARC-Challenge, OpenBookQA, and BoolQ. Fixed-target BPB uses FinePDFs, PG-19, and Proof-Pile. All reported processes completed without failed, OOM, unsupported, missing, or null cells.

Text coverage and aggregation. The 24,000 retrieval evaluations count five models, each on $2 \times 2 \times 8 \times 3 \times 50 = 4,800$ trials. A single stream of concatenated FinePDFs records supplies every real-text haystack; shorter lengths use prefixes of that stream. The same 50 sample indices per task determine cues and five random digits across lengths, depths, suites, and models. These are paired placements and targets rather than independently sampled documents. Synthetic haystacks repeat

the same filler. At each length, overall retrieval gives equal weight to the 12 task-suite-depth cells. Smoke runs are excluded from all reported full-evaluation means.

For primary fixed-target BPB, the evaluator selects the first qualifying documents with at least $262,144 + 256$ GPT-2 tokens: eight FinePDFs, five PG-19, and five Proof-Pile documents. Every context length predicts the same 256-token span immediately after position 262,144, giving 4,608 scored target tokens per model and length. Dataset BPB divides summed target negative log-likelihood in bits by summed UTF-8 target bytes; the reported mean weights the three datasets equally. Additional context changes the conditioning prefix rather than the target. These 18 spans provide a controlled extrapolation diagnostic, not a representative estimate over all documents in the three corpora.

For proportions we report simple Wilson intervals where they aid interpretation. These intervals treat trials as binomial observations and do not exploit the paired design or account for correlations induced by shared templates. The archived retrieval payloads retain aggregate cell counts but the BABILong payloads do not retain per-model correctness vectors, so paired bootstrap or McNemar tests cannot be reconstructed without rerunning or recovering the original predictions.

864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896

E COMPLETE FACTORIAL SWEEP

Table 9 retains all 120 factorial cells; the main paper reports only the recipe transitions needed to justify promotion.

Table 9: Complete Results of the 120-Run Ablation Sweep. All models are 370M parameters, trained on FineWeb-Edu for 1B tokens ($seq_len = 2048$). Evaluated from 2K to 64K context length. Clean NLL is measured on coherent FinePDFs documents; junk NLL is measured on the concatenated validation stream. NLL is in nats/token (lower is better), Needle is greedy per-digit accuracy in % (higher is better), and MFU is training model flops utilization in %.

Type	Reg	Dist	Mem	Win	Val	Clean2K	Clean64K	Junk2K	Junk64K	Ndl2K	Ndl64K	MFU
POLAR	baseline	Off	Off	Off	3.323	2.83	3.60	3.27	5.68	96.3	0.0	40.1
POLAR	baseline	Off	Off	On	3.343	2.83	2.66	3.30	4.74	67.5	0.0	39.6
POLAR	baseline	Off	On	Off	3.169	2.70	1.96	3.11	3.13	91.3	92.5	35.6
POLAR	baseline	Off	On	On	3.174	2.71	2.01	3.12	3.13	51.3	30.0	35.6
POLAR	baseline	On	Off	Off	3.332	2.81	5.56	3.28	6.54	85.0	0.0	33.3
POLAR	baseline	On	Off	On	3.361	2.85	2.89	3.35	4.97	91.3	2.5	34.5
POLAR	baseline	On	On	Off	3.178	2.72	1.98	3.12	3.14	73.8	58.8	31.8
POLAR	baseline	On	On	On	3.182	2.73	1.97	3.12	3.14	56.3	21.3	31.2
POLAR	weak	Off	Off	Off	3.333	2.82	2.86	3.28	5.35	88.8	0.0	39.0
POLAR	weak	Off	Off	On	3.348	2.90	3.76	3.38	5.24	93.8	0.0	37.9
POLAR	weak	Off	On	Off	3.179	2.74	1.95	3.12	3.14	70.0	51.3	36.2
POLAR	weak	Off	On	On	3.187	2.74	1.99	3.13	3.15	41.3	41.3	36.2
POLAR	weak	On	Off	Off	3.335	2.82	3.47	3.28	5.56	91.3	0.0	32.8
POLAR	weak	On	Off	On	3.353	2.82	2.58	3.32	4.81	96.3	1.3	32.8
POLAR	weak	On	On	Off	3.178	2.74	1.95	3.12	3.14	33.8	32.5	30.7
POLAR	weak	On	On	On	3.183	2.74	1.94	3.13	3.15	81.3	78.8	30.7
POLAR	strong	Off	Off	Off	3.338	2.85	3.79	3.28	5.83	98.8	1.3	37.3
POLAR	strong	Off	Off	On	3.360	2.85	2.99	3.32	5.01	91.3	0.0	37.9
POLAR	strong	Off	On	Off	3.172	2.72	1.96	3.11	3.13	80.0	52.5	35.8
POLAR	strong	Off	On	On	3.181	2.72	1.94	3.12	3.15	76.3	68.8	34.6

Continued on next page

Table 9: Complete Results of the 120-Run Ablation Sweep (continued).

Type	Reg	Dist	Mem	Win	Val	Clean2K	Clean64K	Junk2K	Junk64K	Ndl2K	Ndl64K	MFU
POLAR	strong	On	Off	Off	3.337	2.87	3.22	3.28	5.52	83.8	1.3	32.8
POLAR	strong	On	Off	On	3.347	2.82	2.64	3.30	4.88	97.5	16.3	32.2
POLAR	strong	On	On	Off	3.179	2.73	1.94	3.12	3.14	86.3	83.8	30.9
POLAR	strong	On	On	On	3.182	2.72	1.94	3.12	3.14	25.0	23.8	30.9
POLAR	discrete	Off	Off	Off	3.338	2.82	2.80	3.29	5.12	93.8	2.5	37.9
POLAR	discrete	Off	Off	On	3.349	2.87	2.90	3.33	5.25	93.8	6.3	37.9
POLAR	discrete	Off	On	Off	3.184	2.75	2.03	3.12	3.14	76.3	73.8	36.4
POLAR	discrete	Off	On	On	3.178	2.71	1.92	3.12	3.15	31.3	18.8	35.8
POLAR	discrete	On	Off	Off	3.333	2.80	4.13	3.28	6.12	97.5	0.0	32.6
POLAR	discrete	On	Off	On	3.351	2.87	2.72	3.33	4.89	85.0	2.5	33.3
POLAR	discrete	On	On	Off	3.189	2.74	1.96	3.13	3.15	87.5	53.8	31.8
POLAR	discrete	On	On	On	3.192	2.74	2.07	3.13	3.15	21.3	26.3	31.7
POLAR	zipfian	Off	Off	Off	3.338	2.85	4.83	3.28	6.55	96.3	11.3	37.9
POLAR	zipfian	Off	Off	On	3.350	2.83	2.58	3.32	4.63	88.8	3.8	39.4
POLAR	zipfian	Off	On	Off	3.171	2.70	1.92	3.11	3.13	52.5	63.8	36.2
POLAR	zipfian	Off	On	On	3.180	2.75	1.95	3.12	3.15	46.3	40.0	35.2
POLAR	zipfian	On	Off	Off	3.340	2.83	3.77	3.28	5.64	98.8	1.3	33.8
POLAR	zipfian	On	Off	On	3.352	2.85	3.13	3.36	5.23	90.0	0.0	33.5
POLAR	zipfian	On	On	Off	3.177	2.72	1.94	3.12	3.14	91.3	67.5	31.8
POLAR	zipfian	On	On	On	3.192	2.74	2.00	3.13	3.15	40.0	26.3	31.8
NOPE	baseline	Off	Off	Off	3.224	2.76	3.71	3.16	6.68	97.5	1.3	41.5
NOPE	baseline	Off	Off	On	3.219	2.75	2.76	3.17	5.49	92.5	10.0	36.5
NOPE	baseline	Off	On	Off	3.140	2.66	2.34	3.09	3.23	97.5	16.3	36.8
NOPE	baseline	Off	On	On	3.150	2.70	2.14	3.09	3.24	81.3	0.0	34.2
NOPE	baseline	On	Off	Off	3.209	2.74	2.99	3.15	6.13	90.0	18.8	32.5
NOPE	baseline	On	Off	On	3.212	2.77	2.90	3.17	5.65	81.3	0.0	30.3
NOPE	baseline	On	On	Off	3.149	2.67	2.35	3.10	3.28	80.0	16.3	30.0
NOPE	baseline	On	On	On	3.147	2.68	2.09	3.09	3.23	88.8	21.3	28.2
NOPE	weak	Off	Off	Off	3.214	2.77	3.10	3.16	6.34	95.0	0.0	40.9
NOPE	weak	Off	Off	On	3.222	2.75	3.05	3.19	5.81	81.3	0.0	36.1
NOPE	weak	Off	On	Off	3.147	2.68	2.07	3.09	3.19	88.8	21.3	37.5

Continued on next page

930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962

Table 9: Complete Results of the 120-Run Ablation Sweep (continued).

Type	Reg	Dist	Mem	Win	Val	Clean2K	Clean64K	Junk2K	Junk64K	Ndl2K	Ndl64K	MFU
NOPE	weak	Off	On	On	3.142	2.68	2.12	3.08	3.23	82.5	0.0	33.7
NOPE	weak	On	Off	Off	3.215	2.75	4.87	3.16	12.48	95.0	5.0	32.0
NOPE	weak	On	Off	On	3.221	2.73	2.63	3.18	4.84	81.3	0.0	29.3
NOPE	weak	On	On	Off	3.147	2.67	2.48	3.09	3.23	97.5	1.3	30.8
NOPE	weak	On	On	On	3.148	2.68	2.14	3.09	3.24	93.8	1.3	28.4
NOPE	strong	Off	Off	Off	3.231	2.74	4.88	3.18	11.85	90.0	2.5	38.4
NOPE	strong	Off	Off	On	3.214	2.72	2.68	3.17	5.28	82.5	0.0	36.4
NOPE	strong	Off	On	Off	3.144	2.68	2.37	3.09	3.24	96.3	6.3	35.7
NOPE	strong	Off	On	On	3.150	2.70	2.10	3.10	3.25	93.8	7.5	33.8
NOPE	strong	On	Off	Off	3.207	2.72	3.13	3.15	6.20	73.8	0.0	31.4
NOPE	strong	On	Off	On	3.214	2.74	2.49	3.17	4.36	86.3	8.8	28.9
NOPE	strong	On	On	Off	3.146	2.67	2.24	3.09	3.23	83.8	17.5	29.1
NOPE	strong	On	On	On	3.147	2.68	2.16	3.09	3.28	96.3	0.0	27.5
NOPE	discrete	Off	Off	Off	3.202	2.75	3.33	3.15	6.23	85.0	1.3	38.8
NOPE	discrete	Off	Off	On	3.234	2.79	3.11	3.19	5.47	70.0	18.8	37.0
NOPE	discrete	Off	On	Off	3.145	2.68	2.14	3.09	3.27	98.8	16.3	36.0
NOPE	discrete	Off	On	On	3.145	2.68	2.10	3.09	3.25	95.0	3.8	34.3
NOPE	discrete	On	Off	Off	3.212	2.88	3.04	3.16	6.88	42.5	0.0	31.2
NOPE	discrete	On	Off	On	3.211	2.78	2.73	3.16	5.28	82.5	1.3	29.1
NOPE	discrete	On	On	Off	3.145	2.68	2.27	3.08	3.26	91.3	11.3	30.3
NOPE	discrete	On	On	On	3.149	2.68	2.17	3.08	3.32	83.8	0.0	28.6
NOPE	zipfian	Off	Off	Off	3.209	2.71	3.18	3.15	6.40	90.0	0.0	40.8
NOPE	zipfian	Off	Off	On	3.219	2.75	2.71	3.17	5.44	77.5	1.3	36.1
NOPE	zipfian	Off	On	Off	3.141	2.67	2.55	3.08	3.38	96.3	0.0	38.0
NOPE	zipfian	Off	On	On	3.150	2.69	2.11	3.09	3.25	90.0	3.8	33.7
NOPE	zipfian	On	Off	Off	3.207	2.77	3.32	3.15	7.09	82.5	0.0	31.3
NOPE	zipfian	On	Off	On	3.227	2.79	2.91	3.18	5.94	93.8	11.3	30.2
NOPE	zipfian	On	On	Off	3.146	2.68	2.39	3.10	3.25	96.3	21.3	29.6
NOPE	zipfian	On	On	On	3.148	2.68	2.17	3.09	3.28	90.0	6.3	28.7
ROPE	baseline	Off	Off	Off	3.159	2.85	3.36	3.09	5.21	42.5	0.0	40.4
ROPE	baseline	Off	Off	On	3.163	2.86	3.79	3.17	5.86	15.0	0.0	37.2

Continued on next page

Table 9: Complete Results of the 120-Run Ablation Sweep (continued).

Type	Reg	Dist	Mem	Win	Val	Clean2K	Clean64K	Junk2K	Junk64K	Ndl2K	Ndl64K	MFU
ROPE	baseline	Off	On	Off	3.145	2.81	2.86	3.09	3.57	73.8	0.0	37.1
ROPE	baseline	Off	On	On	3.151	2.80	2.95	3.12	3.72	13.8	0.0	34.4
ROPE	baseline	On	Off	Off	3.163	3.00	3.47	3.10	5.46	38.8	0.0	33.2
ROPE	baseline	On	Off	On	3.167	2.88	4.06	3.16	5.92	0.0	0.0	30.2
ROPE	baseline	On	On	Off	3.153	2.73	2.45	3.09	3.57	75.0	0.0	31.4
ROPE	baseline	On	On	On	3.154	2.83	2.86	3.12	3.72	32.5	0.0	29.4
ROPE	weak	Off	Off	Off	3.158	2.88	3.46	3.10	5.36	46.3	0.0	41.4
ROPE	weak	Off	Off	On	3.163	2.84	3.96	3.18	5.93	12.5	0.0	37.8
ROPE	weak	Off	On	Off	3.149	2.71	2.46	3.09	3.56	37.5	0.0	36.5
ROPE	weak	Off	On	On	3.156	2.85	2.77	3.12	3.60	23.8	0.0	34.9
ROPE	weak	On	Off	Off	3.155	2.83	3.49	3.10	5.27	38.8	0.0	32.3
ROPE	weak	On	Off	On	3.157	3.05	4.12	3.18	5.86	8.8	0.0	30.5
ROPE	weak	On	On	Off	3.149	2.77	2.76	3.09	3.56	57.5	0.0	31.1
ROPE	weak	On	On	On	3.149	2.77	3.13	3.12	3.80	17.5	0.0	28.0
ROPE	strong	Off	Off	Off	3.165	2.78	3.34	3.10	5.27	30.0	0.0	39.0
ROPE	strong	Off	Off	On	3.169	2.90	3.68	3.19	5.80	31.3	0.0	36.3
ROPE	strong	Off	On	Off	3.158	2.85	2.63	3.10	3.50	51.3	0.0	35.9
ROPE	strong	Off	On	On	3.162	2.86	3.08	3.14	3.76	10.0	0.0	33.2
ROPE	strong	On	Off	Off	3.161	2.85	3.58	3.10	5.28	47.5	0.0	31.1
ROPE	strong	On	Off	On	3.165	2.85	4.03	3.19	5.90	2.5	0.0	29.2
ROPE	strong	On	On	Off	3.155	2.80	2.70	3.10	3.62	61.3	0.0	29.3
ROPE	strong	On	On	On	3.162	2.85	3.08	3.13	3.74	2.5	0.0	28.4
ROPE	discrete	Off	Off	Off	3.165	2.91	3.43	3.11	5.42	22.5	0.0	39.7
ROPE	discrete	Off	Off	On	3.158	2.84	4.05	3.15	5.90	3.8	0.0	36.9
ROPE	discrete	Off	On	Off	3.159	2.83	2.67	3.11	3.56	47.5	0.0	36.3
ROPE	discrete	Off	On	On	3.161	2.79	2.71	3.13	3.74	41.3	0.0	33.7
ROPE	discrete	On	Off	Off	3.171	2.79	3.45	3.12	5.36	51.3	0.0	31.5
ROPE	discrete	On	Off	On	3.164	2.80	3.74	3.16	5.86	21.3	0.0	30.0
ROPE	discrete	On	On	Off	3.162	2.74	2.47	3.11	3.53	50.0	0.0	30.9
ROPE	discrete	On	On	On	3.154	2.79	2.68	3.11	3.67	18.8	0.0	28.9
ROPE	zipfian	Off	Off	Off	3.163	2.82	3.25	3.10	5.22	62.5	0.0	39.8

Continued on next page

996
997
998
999
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028

Table 9: Complete Results of the 120-Run Ablation Sweep (continued).

Type	Reg	Dist	Mem	Win	Val	Clean2K	Clean64K	Junk2K	Junk64K	Ndl2K	Ndl64K	MFU
ROPE	zipfian	Off	Off	On	3.161	2.94	3.94	3.18	5.81	16.3	0.0	37.1
ROPE	zipfian	Off	On	Off	3.149	2.78	2.70	3.09	3.63	46.3	0.0	36.7
ROPE	zipfian	Off	On	On	3.151	2.80	2.63	3.11	3.68	2.5	0.0	33.9
ROPE	zipfian	On	Off	Off	3.161	2.78	3.38	3.10	5.32	47.5	0.0	32.6
ROPE	zipfian	On	Off	On	3.160	3.01	4.11	3.17	5.81	17.5	0.0	30.2
ROPE	zipfian	On	On	Off	3.149	2.73	2.39	3.09	3.55	62.5	0.0	29.8
ROPE	zipfian	On	On	On	3.152	2.79	2.99	3.12	3.78	26.3	0.0	28.5

F SEED-PAIRED REPLICATION

We trained four fresh 378M-parameter checkpoints: paired NoPE and Polar runs at seeds 202701 and 202702. Every run used 18,722 steps of 524,288 tokens (9.816B tokens), sequence length 2,048, the L40S configuration with microbatch size four, and the same data order within each seed pair. Final validation NLL differed by architecture but was stable across seeds (Table 10). Checkpoints and evaluation datasets were pinned to immutable Hub revisions.

The follow-up was intentionally restricted to the two diagnostics relevant to the retention hypothesis: teacher-forced retrieval and fixed-target BPB. We did not run downstream tasks or BABILong. For each checkpoint we compared the untouched model with the checkpoint-local 256-token cap on its largest parameter-only zero-input retention head. Retrieval used the same synthetic/FinePDFs, passkey/NIAH, length, depth, sample, target, and seed grid as the primary evaluation: 38,400 trials across the eight model-condition combinations. Fixed-target BPB used eight qualifying FinePDFs documents and five qualifying documents each from PG-19 and Proof-Pile; document hashes were identical across models and conditions. All 24 jobs completed without an OOM or failed cell.

Seed	Architecture	Final val. NLL	Max $H_{1/2}$	256K BPB	256K retrieval tok./exact
202701	NoPE	2.81164	135k (B2/H2)	3.898 $\rightarrow$ 1.506	4.83/0.00 $\rightarrow$ 0.73/0.00
202701	Polar	2.85661	272 (B10/H7)	1.644 $\rightarrow$ 1.644	16.97/0.00 $\rightarrow$ 16.53/0.00
202702	NoPE	2.81474	5.78M (B2/H3)	5.256 $\rightarrow$ 1.490	12.13/0.00 $\rightarrow$ 7.07/0.00
202702	Polar	2.85417	231 (B6/H4)	1.760 $\rightarrow$ 1.762	20.90/3.83 $\rightarrow$ 21.13/3.83

Table 10: Seed-level replication endpoints. Each metric reports untouched $\rightarrow$ capped evaluation. Retrieval averages passkey and NIAH, three depths, and synthetic and real-text suites; values are percentages. $H_{1/2}$ is the maximum parameter-only zero-input operating point, with zero-indexed block/head coordinates.

Architecture	Condition	Metric	2K	4K	8K	16K	32K	64K	128K	256K
NoPE	untouched	Retrieval token	99.00	72.50	51.62	47.62	42.42	33.85	19.42	8.48
NoPE	capped	Retrieval token	98.55	85.27	82.52	56.53	34.05	17.77	8.50	3.90
Polar	untouched	Retrieval token	98.65	95.77	97.52	89.17	67.32	51.30	27.88	18.93
Polar	capped	Retrieval token	98.70	95.78	97.77	89.48	67.40	52.25	28.22	18.83
NoPE	untouched	Retrieval exact	95.00	61.33	43.25	39.25	24.50	9.08	0.75	0.00
NoPE	capped	Retrieval exact	92.92	62.58	55.25	19.92	5.08	0.08	0.00	0.00
Polar	untouched	Retrieval exact	93.58	84.25	88.67	70.08	18.33	12.42	4.83	1.92
Polar	capped	Retrieval exact	93.83	84.42	89.50	70.17	19.08	14.33	5.17	1.92
NoPE	untouched	Mean BPB	1.379	1.442	2.016	2.590	3.346	3.619	3.902	4.577
NoPE	capped	Mean BPB	1.397	1.397	1.416	1.420	1.445	1.435	1.459	1.498
Polar	untouched	Mean BPB	1.432	1.424	1.448	1.458	1.522	1.557	1.617	1.702
Polar	capped	Mean BPB	1.435	1.426	1.449	1.459	1.523	1.558	1.618	1.703

Table 11: Complete replication curves, descriptively averaged over the two seeds. Retrieval metrics are percentages and give equal weight to the 12 task-suite-depth cells at each length. BPB gives equal weight to FinePDFs, PG-19, and Proof-Pile. These two seeds do not support a population uncertainty estimate.

Both NoPE checkpoints developed much longer zero-input horizons than either Polar checkpoint. Their mean BPB growth from 2K to 256K was $3.32\times$ when untouched and $1.07\times$ when capped; the cap repaired the likelihood collapse in both seeds. It did not rescue retrieval: seed-averaged NoPE token accuracy fell from 33.85% to 17.77% at 64K and from 8.48% to 3.90% at 256K. Conversely, Polar had $1.19\times$ BPB growth in both conditions and nearly overlapping retrieval curves because its maxima were already near the selected ceiling. At 256K, every condition had zero exact real-text retrieval; Polar’s 1.92% overall exact mean came from the synthetic suite in seed 202702. The replication therefore supports excessive retention as a checkpoint-specific likelihood-collapse mech-

Model	Haystack	2K	4K	8K	16K	32K	64K	128K	256K
Polar	Synthetic	99.1	98.4	97.1	97.0	94.9	92.3	80.3	68.4
	FinePDFs	98.6	86.3	77.9	49.1	40.7	16.3	3.2	0.3
NoPE	Synthetic	99.1	98.7	97.7	94.9	70.3	26.5	10.2	1.2
	FinePDFs	99.3	52.9	0.0	0.0	0.0	0.0	0.0	0.0
RoPE	Synthetic	81.7	44.9	34.5	0.8	6.8	0.1	0.0	0.0
	FinePDFs	94.8	4.7	0.2	0.0	0.0	0.0	0.0	0.0
Atma-Raven-Titans	Synthetic	63.0	60.7	53.8	46.9	41.7	35.8	25.9	19.3
	FinePDFs	23.2	2.6	0.7	0.2	1.4	0.5	1.6	0.7
Raven Native	Synthetic	60.1	57.7	52.7	48.3	44.3	40.9	38.9	35.1
	FinePDFs	22.1	15.4	8.2	3.3	5.3	2.9	3.7	5.4

Table 12: Untouched teacher-forced target-token retrieval accuracy (%). Each entry averages two tasks and three depths, with 50 trials per cell. Only full evaluations enter these means; smoke runs are excluded.

anism, especially for NoPE, while refuting a general expectation that the cap improves retrieval or that every Polar checkpoint develops a comparable outlier.

G RETRIEVAL BREAKDOWN

Retrieval token accuracy by context length and needle depth




Figure 5: Retrieval depth sweeps for the three matched attention variants before and after the one-head retention cap, followed by untouched Raven Native and Atma-Raven-Titans reference panels. Each cell is teacher-forced token accuracy averaged over passkey and NIAH and over synthetic and FinePDFs haystacks. The Raven rows are separately optimized comparisons, not clamped ablations.

Figure 5 shows variation by needle depth across all five models. Tables 12 and 13 separate partial target-token signal from exact five-token recall. Every tested model eventually reaches 0.0% exact retrieval on real-text haystacks under unadapted length extrapolation:

Model	Haystack	2K	4K	8K	16K	32K	64K	128K	256K
Polar	Synthetic	96.3	92.3	86.3	86.0	76.0	65.7	26.3	18.0
	FinePDFs	93.0	67.3	62.7	12.0	0.0	0.0	0.0	0.0
NoPE	Synthetic	95.3	93.7	88.7	78.0	23.7	4.7	0.7	0.0
	FinePDFs	96.7	18.0	0.0	0.0	0.0	0.0	0.0	0.0
RoPE	Synthetic	42.3	4.3	0.7	0.0	0.0	0.0	0.0	0.0
	FinePDFs	77.3	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Atma-Raven-Titans	Synthetic	8.0	5.7	3.3	2.0	0.7	0.3	0.0	0.0
	FinePDFs	4.7	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Raven Native	Synthetic	2.0	1.0	0.3	0.0	0.0	0.0	0.0	0.0
	FinePDFs	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Table 13: Untouched teacher-forced exact five-token retrieval accuracy (%). Each entry averages two tasks and three depths, with 50 trials per cell. Only full evaluations enter these means; smoke runs are excluded.

- **Raven Native and Atma-Raven-Titans:** Raven Native scores 0.0% on FinePDFs at every tested length and no more than 2.0% on synthetic haystacks. Atma-Raven-Titans scores 4.7% at 2K on FinePDFs and 0.0% from 4K onward.
- **RoPE:** Exact FinePDFs retrieval falls from 77.3% at 2K to 0.0% at 4K and stays there through 256K.
- **NoPE:** Exact FinePDFs retrieval falls from 96.7% at 2K to 18.0% at 4K and 0.0% from 8K onward.
- **Polar Attention:** Untouched Polar reaches 93.0%, 67.3%, 62.7%, and 12.0% on FinePDFs at 2K, 4K, 8K, and 16K. On synthetic haystacks, it retains 65.7% exact retrieval at 64K and 18.0% at 256K, with a 95% Wilson interval of 14.1–22.7% at 256K. The other models score 0.0% there.

Thus, the loss of exact natural-text retrieval at long context is not unique to ATMA. Polar delays this loss relative to the tested baselines but does not eliminate it.

G.1 OPEN-WEIGHT CONTEXT

Table 14 reports the ten archived open-weight controls. They differ in training data, tokenizer, architecture, parameter count, and native context length; several were trained at lengths beyond 2K. They are therefore practical context, not matched length-extrapolation baselines and not evidence of a Polar advantage over the corresponding model families.

Open-weight model	NLL 2K	NLL 64K	Needle 2K	Needle 64K
SmolLM2-360M	1.99	6.04	100.0	8.1
LFM2.5-230M	3.01	6.69	38.1	19.4
LFM2.5-350M	3.18	6.63	45.6	46.3
Qwen2.5-0.5B	2.03	1.53	100.0	38.1
Qwen3-0.6B-Base	1.95	1.40	100.0	95.0
Qwen3.5-0.8B-Base	1.94	1.33	100.0	100.0
Gemma-3-270M	2.48	1.90	96.2	24.4
Granite-4.0-350M	2.52	3.02	40.6	8.8
Granite-4.0-H-350M	2.22	2.06	75.0	54.4
Falcon-H1-0.5B	1.99	2.54	92.5	19.4

Table 14: Open-weight controls evaluated through 64K. NLL is fixed-target FinePDFs nats/token; Needle is teacher-forced target-token accuracy over 16 trials. These checkpoints are not matched for data, tokenizer, native context, or training budget.

Group	Model	LAMBADA	HellaSwag	PIQA	WinoGrande	ARC-E	ARC-C	OBQA	BoolQ	Mean
Matched	NoPE	30.16	37.88	66.27	51.93	50.53	32.44	31.60	60.37	45.15
Matched	RoPE	30.25	37.19	66.76	50.67	49.12	29.43	32.40	60.95	44.60
Matched	Polar	28.47	36.19	66.87	51.62	49.12	26.09	31.80	56.21	43.29
Raven/AdamW	Atma-Raven-Titans	26.00	34.12	65.56	52.49	48.07	27.76	29.00	61.41	43.05
Raven/AdamW	Raven Native	27.48	33.56	64.69	51.70	49.12	27.09	28.80	61.93	43.05

Table 15: Task-level zero-shot accuracy (%) on the eight short-context controls (22,939 examples per model). LAMBADA, WinoGrande, and BoolQ use standard accuracy; HellaSwag, PIQA, ARC-Easy, ARC-Challenge, and OpenBookQA use length-normalized accuracy. Mean is the unweighted average of the eight displayed metrics; bold is best within each comparison group.




Figure 6: Single-sequence serving on one L40S. Attention decode scales with KV-cache length; recurrent Raven uses fixed state.

H SHORT-CONTEXT QUALITY AND SYSTEMS

Table 15 gives the exact per-task controls shown in main-text Figure 3. Figure 6 shows the serving-length trend, and Table 16 summarizes the quality–systems trade-off.

I SYSTEMS IMPLEMENTATION AND OPTIMIZATION BREAKDOWN

The inference stack uses separate kernels for full-context prefill, paged autoregressive decode, and recurrent-state updates. Table 17 distinguishes what each optimization changes. In particular, streaming removes quadratic score storage but not quadratic prefill arithmetic, while the recurrent Raven path changes the state representation itself and therefore has different scaling.

Memory accounting. Engine allocation includes weights, runtime buffers, and any preallocated KV pool. The attention runner sizes its pool against a 90% device-memory budget, whereas Raven allocates recurrent state without a KV pool. The roughly 39.4 GiB attention values therefore reflect the configured engine’s allocation policy and do not measure the minimum memory required for one 128K sequence. With four global layers, two KV heads, head dimension 128, and BF16 storage, the occupied key/value tensors for that sequence require $2 \times 4 \times 131,072 \times 2 \times 128 \times 2$ bytes, or 0.5 GiB, before weights, activations, recurrent state, and other buffers. This is an analytical KV-size calculation, not a new peak-memory measurement. The recurrent-state and KV-length complexity comparison remains separate from these allocated-pool measurements.

Group	Model	Base mean (%)	128K prefill (s)	Decode (ms/token)	Engine alloc. (GiB)	Train MFU (%, 2K)
Matched	NoPE	45.1	1.47	16.4	39.40	42.77
Matched	Polar	43.3	1.62	16.3	39.41	40.44
Matched	RoPE	44.6	1.47	16.5	39.41	42.45
Raven/AdamW	Atma-Raven-Titans	43.1	1.00	2.27	8.62	30.92
Raven/AdamW	Raven Native	43.0	1.36	3.27	6.46	21.53

Table 16: Short-context quality and serving trade-offs. Serving uses one sequence and one timing sample. Engine allocation includes the attention backend’s preallocated KV pool; it is not occupied per-sequence cache size. See the memory-accounting discussion below.

Path	Optimization	Operational effect
Polar prefill	Block-streamed K, V tiles with on-line (M, L, Q_2, S) accumulation	Avoids materializing $T \times T$ scores; retains full-context $O(T^2)$ arithmetic.
Polar decode	Paged KV-cache kernel grouping the four query heads that share each KV head	Reuses KV loads across the GQA group; decode remains $O(T)$ in cache length.
Gated-delta step	Fused in-place decay, prediction, delta write, and readout	Avoids materializing per-step intermediate states and preserves $O(1)$ recurrent state size.
Ragged prefill	Packed prompts with grouped compatible tile shapes	Reduces padding work on heterogeneous prompt batches.
Compiler boundary	Recurrent scan and custom kernels exposed as opaque operations	Prevents compiler graph breaks around the stateful update while retaining explicit numerical tests.

Table 17: Implementation-level optimization breakdown. Complexity statements describe the sequence-mixing path and exclude shared MLP and embedding costs.

1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331

I.1 SINGLE-SEQUENCE SCALING

Table 18 gives the underlying single-sequence measurements. On one NVIDIA L40S, attention decode increases from approximately 2.3 ms/token at 2K to 16.3–16.5 ms/token at 128K because every step reads a longer KV cache. The recurrent models remain nearly flat because they update fixed-size state. Prefill is less monotonic at short lengths because launch overhead and kernel occupancy dominate before the matrix operations reach their efficient regime. These timings use batch size 1, 32 generated tokens, and one recorded timing sample per cell. They are descriptive measurements, not confidence intervals.

Table 18: Detailed prefill and decode scaling from 2K to 128K tokens on one NVIDIA L40S GPU (batch size 1, 32 generated tokens).

Model	Metric	2K	4K	8K	16K	32K	64K	128K
Polar	Prefill latency (s)	0.131	0.031	0.051	0.101	0.248	0.601	1.619
	Prefill throughput (tok/s)	15,640	133,884	160,145	161,551	132,091	109,050	80,959
	Decode latency (ms/tok)	2.27	2.49	2.88	3.70	5.66	9.22	16.30
	Decode throughput (tok/s)	441.2	402.3	346.6	270.1	176.8	108.5	61.3
NoPE	Prefill latency (s)	0.157	0.028	0.049	0.097	0.230	0.576	1.473
	Prefill throughput (tok/s)	13,070	146,568	166,504	168,510	142,319	113,810	89,013
	Decode latency (ms/tok)	2.22	2.45	2.83	3.68	5.62	9.22	16.40
	Decode throughput (tok/s)	450.1	408.8	353.0	271.7	177.8	108.4	61.0
RoPE	Prefill latency (s)	0.160	0.028	0.049	0.096	0.232	0.566	1.472
	Prefill throughput (tok/s)	12,761	147,089	168,287	170,936	141,249	115,786	89,015
	Decode latency (ms/tok)	2.34	2.53	2.95	3.79	5.73	9.33	16.51
	Decode throughput (tok/s)	427.4	395.3	338.7	263.7	174.5	107.2	60.6
Atma-Raven-Titans	Prefill latency (s)	0.198	0.041	0.059	0.111	0.247	0.499	0.999
	Prefill throughput (tok/s)	10,334	100,989	137,903	147,434	132,916	131,464	131,195
	Decode latency (ms/tok)	2.24	2.25	2.24	2.24	2.26	2.25	2.27
	Decode throughput (tok/s)	446.6	445.1	446.0	447.1	442.6	444.8	440.3
Raven Native	Prefill latency (s)	0.214	0.054	0.070	0.139	0.320	0.672	1.357
	Prefill throughput (tok/s)	9,558	75,738	116,378	117,605	102,354	97,527	96,559
	Decode latency (ms/tok)	3.26	3.25	3.24	3.24	3.25	3.25	3.27
	Decode throughput (tok/s)	306.4	307.5	308.2	308.3	307.6	307.5	306.1

I.2 PRODUCTION-FRAMEWORK STRESS COMPARISON

To test whether the custom kernels introduce an obvious systems bottleneck at a larger shape, we configured ATMA’s inference engine with a 9.2B-parameter stress shape and compared it with vLLM 0.25.0 serving Qwen3.5-9B on the same L40S in BF16 (Kwon et al., 2023). The models are not weight- or architecture-matched, so Table 19 is a backend stress comparison rather than a model-quality or universal serving claim. ATMA trails on short homogeneous prefill, is similar on long heterogeneous prefill, and leads the recorded high-batch decode cells.

Phase	Batch/workload	ATMA engine	vLLM/Qwen3.5-9B	Ratio
Prefill	Homogeneous $B = 8, T = 512$ (4,096 tok)	11,947 tok/s	12,772 tok/s	$0.94\times$
	Heterogeneous short-heavy (816 tok)	5,752 tok/s	8,002 tok/s	$0.72\times$
	Heterogeneous mixed (4,384 tok)	12,779 tok/s	11,996 tok/s	$1.07\times$
	Heterogeneous long-tail (8,192 tok)	12,251 tok/s	11,721 tok/s	$1.05\times$
Decode, context 512	Batch 128	2,939 tok/s	2,247 tok/s	$1.31\times$
	Batch 256	4,501 tok/s	2,921 tok/s	$1.54\times$
	Batch 512	5,919 tok/s	2,558 tok/s	$2.31\times$
	Heterogeneous $B = 128$ (context 64–1536)	2,950 tok/s	2,403 tok/s	$1.23\times$

Table 19: Recorded systems-scale throughput for ATMA’s 9.2B stress shape and vLLM 0.25.0 serving Qwen3.5-9B on one L40S (BF16, greedy sampling). Ratio is ATMA/vLLM.

J RETENTION-HORIZON DIAGNOSTIC

This section records the post-hoc sequence from parameter inspection to a multi-clamp sweep and broad benchmark re-evaluation. It analyzes released checkpoints rather than a new trained architecture. The primary checkpoint comparisons in Section 5 use untouched models. Every cap below is opt-in, applied only at inference, and removed after each condition.

The diagnostic follows three systematic stages:

1. **Parameter-only operating point inspection:** Identifying unusually large learned retention gate biases (b_γ) without running the model.
2. **Causal multi-clamping sweep:** Evaluating multiple clamp values (local percentiles p_{90} , p_{99} and absolute half-life ceilings $h_{l:256}$, $h_{l:512}$) across sequence lengths 2K–256K on clean text loss and needle retrieval.
3. **Independent benchmark re-evaluation:** Re-benchmarking the selected intervention across short-context downstream tasks, retrieval (synthetic and FinePDFs real-text haystacks, passkey and NIAH), fixed-target long-document likelihood (FinePDFs, PG-19, Proof-Pile), and BABILong reasoning.

J.1 PARAMETER-ONLY OPERATING POINTS

The memory gate has one learned row per query/memory head, rather than one scalar per attention KV head. For each row we decompose the logit as

$$z_{\gamma,t} = W_\gamma x_t + b_\gamma, \quad \gamma_0 = \sigma(b_\gamma), \quad H_{1/2}(\gamma_0) = \frac{\log(1/2)}{\log \gamma_0}. \quad (11)$$

Here $x = 0$ defines a reproducible learned operating point. It is not the empirical mean runtime gate: every reported outlier has a nonzero W_γ row, so activations can shorten or lengthen its instantaneous horizon.

Table 20 reports the parameter operating points across all base checkpoints. The promoted NoPE and Polar checkpoints trained on L40S contain isolated Block 2 heads with learned half-lives of 21.01M and 3.07M tokens. RoPE has a maximum $H_{1/2}$ of 682.2 tokens, while Atma-Raven-Titans ranges from 30.5 to 41.4 tokens. We therefore do not cap Atma-Raven-Titans.

Scanning the separately fine-tuned BABILong checkpoints gives 20.86M tokens for NoPE, 3.05M for Polar, and 680 for RoPE. The outlier therefore remains nearly unchanged after adaptation. Each




Figure 7: Complete retention-intervention curves for the primary checkpoints. Dashed open curves are untouched; solid curves apply the one-head cap. Dotted filled Raven curves are untouched references. The broad re-evaluation uses pinned archived baselines; same-process sweeps and fresh paired replications are reported separately.

BABILong cap is still selected from the final fine-tuned checkpoint rather than copied from the base model.

J.2 PAIRED MULTI-CLAMP SWEEP

For a target maximum half-life $H_{\max}$, the diagnostic clips the final activation-conditioned gate logit,

$$z'_{\gamma,t} = \min(z_{\gamma,t}, \logit(2^{-1/H_{\max}})), \quad \gamma'_t = \sigma(z'_{\gamma,t}). \quad (12)$$

For $h1:256$, $H_{\max} = 256 \implies \gamma'_t \leq 0.997296$. Only the single layer/head with the largest parameter-only b_{γ} is targeted. This upper bound leaves smaller token-conditioned gates unchanged, modifies neither checkpoint tensors nor attention, and is removed after the condition.

Table 21 reports the complete paired sweep across conditions (baseline, p90, p99, h1:256, h1:512) and lengths (2K, 16K, 64K, 256K) within a single process sharing model instance, documents, and seeds:

- **Why h1:256 is selected:** It gives the largest reduction in 256K needle cross-entropy, from 12.20 to 6.05 nats on NoPE and from 2.39 to 1.71 nats on Polar. It also keeps 2K clean-loss drift within the 0.05-nat guardrail, at +0.022 nats on NoPE and +0.010 on Polar.
- **Why p99 fails on saturated checkpoints:** On NoPE L40S, the 99th-percentile logit (13.46) still gives a half-life of about 484k tokens. At 256K, clean loss remains at 12.09 nats and needle cross-entropy at 12.15 nats.
- **Why RoPE is an empirical negative control:** RoPE has no comparable outlier. Capping it at h1:256 gives no retrieval benefit at 256K (0.0% accuracy) and reduces short-context needle accuracy from 82.5% to 50.0%. The intervention is therefore not a general benchmark boost.

J.3 INDEPENDENT BENCHMARK RE-EVALUATION

After the pilot sweep, we evaluate the capped promoted checkpoints on all eight short-context downstream tasks, the full synthetic and FinePDFs retrieval matrix, dataset-level long-document likelihood, and BABILong QA1–QA10. Base-model and BABILong-fine-tuned checkpoints are pinned to their archived Hugging Face revisions. Direct checkpoint-exact scoring is used throughout. All 15 model-suite jobs completed without an out-of-memory failure.

Tables 20–29 report the full numerical arrays:

- **Downstream Controls (Table 22):** Unweighted 2K task means change by at most 0.06 points across all three models, showing small changes in the aggregate downstream mean; this does not imply preservation of every short-context retrieval metric.
- **Retrieval curves (Tables 23, 24, 25, 26):** Polar gains +12.70 points in target-token accuracy and +7.33 points in exact five-token accuracy at 256K. Its 8K FinePDFs exact accuracy rises from 62.7% to 91.3%, and capped exact recall remains nonzero through 128K. NoPE remains at 0.0% exact real-text retrieval from 16K to 256K, and RoPE remains at 0.0% at every extrapolation length.
- **BABILong reasoning (Table 27):** The cap raises NoPE to 35–57% from 8K to 256K, compared with 39/21/6% untouched at 8/16/32K and 0% from 64K onward, and raises Polar from 28% to 42% at 256K. RoPE decreases from 14% to 11%.
- **Long-document likelihood (Tables 28, 29):** At 256K, the cap reduces NoPE BPB from 11.666 to 1.180 on FinePDFs, from 4.902 to 1.224 on PG-19, and from 7.723 to 2.382 on Proof-Pile. Polar’s mean degradation drops from $1.26\times$ to $1.03\times$.

Checkpoint	Role	Outlier Block/Head	Learned b_γ	Logit z_0	γ_0	Max $H_{1/2}$ (tok)	Median $H_{1/2}$	Min $H_{1/2}$	$\ W_\gamma\ _2$
NoPE (L4, mbs4)	Earlier L4 baseline	Block 2, H7	+1.687	5.587	0.9962667	185.3	29.6	3.9	3.36
NoPE (L40S, mbs4)	Microbatch ablation	Block 2, H2	+10.607	14.507	0.9999995	1.38M	21.8	4.9	3.49
NoPE (L40S, mbs16)	Promoted matched	Block 2, H5	+13.327	17.227	1.0000000	21.01M	26.1	3.5	3.33
Polar (L40S, mbs16)	Promoted matched	Block 2, H6	+11.403	15.303	0.9999998	3.07M	34.3	5.2	2.76
RoPE (L40S, mbs16)	Matched negative control	Block 2, H1	+2.991	6.891	0.9989845	682.2	26.4	8.8	2.19
Atma-Raven-Titans	Recurrent reference	Block 2, H3	+0.182	4.082	0.9834064	41.4	34.9	30.5	1.94

Table 20: Learned zero-input retention operating points across base checkpoints. Isolated Block 2 heads in L40S NoPE and Polar have half-lives of $> 3\text{M}$ – 21M tokens. The maxima are 682 tokens for RoPE and 41 tokens for Atma-Raven-Titans.

Model	Condition	Clean NLL (nats)				Needle Cross-Entropy (nats)				Needle Acc. (%)	
		2K	16K	64K	256K	2K	16K	64K	256K	2K	256K
NoPE (L40S, mbs16)	Baseline (none)	2.321	2.791	7.177	12.003	0.103	6.105	10.800	12.198	100.0	0.0
	p90 quantile	2.359	1.442	1.263	1.133	0.428	2.511	5.718	6.060	100.0	0.0
	p99 quantile	2.321	2.764	6.955	12.094	0.103	6.016	10.773	12.148	100.0	0.0
	hl:256 (cap) (sel.)	2.343	1.433	1.263	1.138	0.245	2.024	5.687	6.051	100.0	0.0
	hl:512 (cap)	2.333	1.437	1.279	1.162	0.163	1.799	5.651	6.076	100.0	0.0
Polar (L40S, mbs16)	Baseline (none)	2.341	1.423	1.300	1.252	0.262	0.509	1.571	2.387	100.0	50.0
	p90 quantile	2.353	1.399	1.201	1.057	0.282	0.167	0.369	1.713	97.5	65.0
	p99 quantile	2.342	1.422	1.295	1.237	0.264	0.478	1.529	2.366	100.0	55.0
	hl:256 (cap) (sel.)	2.351	1.399	1.201	1.057	0.283	0.165	0.400	1.707	100.0	65.0
	hl:512 (cap)	2.349	1.398	1.201	1.059	0.285	0.163	0.506	1.806	100.0	65.0
RoPE (L40S, mbs16)	Baseline (none)	2.343	2.145	2.581	2.948	0.966	4.809	5.975	6.157	82.5	0.0
	p90 quantile	2.365	2.176	2.488	2.745	3.345	5.053	5.870	5.964	17.5	0.0
	p99 quantile	2.346	2.136	2.478	2.758	1.887	4.919	5.906	6.033	55.0	0.0
	hl:256 (cap)	2.347	2.139	2.477	2.754	2.074	4.936	5.897	6.019	50.0	0.0
	hl:512 (cap)	2.345	2.132	2.485	2.774	1.499	4.886	5.922	6.073	70.0	0.0
NoPE (L40S, mbs4)	Baseline (none)	2.291	5.542	8.425	9.589	0.098	8.093	9.428	10.446	100.0	0.0
	p90 quantile	2.319	1.406	1.250	1.122	0.222	1.283	4.988	5.888	100.0	0.0
	p99 quantile	2.291	5.489	8.168	9.236	0.097	8.066	9.382	10.450	100.0	0.0
	hl:256 (cap)	2.304	1.394	1.274	1.181	0.161	1.167	5.127	6.108	100.0	0.0
	hl:512 (cap)	2.298	1.534	1.664	1.549	0.127	2.248	5.353	6.151	100.0	0.0
NoPE (L4, mbs4)	Baseline (none)	2.304	2.546	3.010	3.615	0.163	1.326	2.649	3.730	97.5	15.0
	p90 quantile	2.307	1.828	2.046	2.333	0.181	1.929	2.694	4.447	100.0	5.0
	p99 quantile	2.306	1.934	2.210	2.613	0.177	1.899	2.699	4.435	97.5	10.0
	hl:256 (cap)	2.305	2.016	2.334	2.776	0.174	1.843	2.706	4.357	97.5	10.0
	hl:512 (cap)	2.304	2.181	2.565	3.030	0.171	1.699	2.719	4.142	97.5	10.0

Table 21: Complete paired clamp sweep across quantile and half-life caps from 2K to 256K. Capping the single outlier head at hl:256 improves NoPE and Polar extrapolation while meeting the 2K short-context guardrail. The p99 setting leaves a long retention horizon, while the RoPE cap reduces short-range needle accuracy from 82.5% to 50.0% without improving long-range retrieval.

Model	LAMBADA	HellaSwag	PIQA	WinoGrande	ARC-E	ARC-C	OBQA	BoolQ	Mean
NoPE	30.16 → 30.22	37.88 → 37.88	66.27 → 66.21	51.93 → 51.38	50.53 → 50.70	32.44 → 32.44	31.60 → 31.60	60.37 → 60.34	45.15 → 45.10
Polar	28.47 → 28.22	36.19 → 36.20	66.87 → 67.30	51.62 → 51.38	49.12 → 48.60	26.09 → 26.42	31.80 → 32.00	56.21 → 55.87	43.29 → 43.25
RoPE	30.25 → 30.16	37.19 → 37.17	66.76 → 66.97	50.67 → 51.07	49.12 → 48.95	29.43 → 29.77	32.40 → 32.20	60.95 → 61.01	44.60 → 44.66

Table 22: Complete short-context control, untouched → capped (accuracy, %). The unweighted eight-task mean changes by less than 0.06 percentage points for each model; individual task changes can be larger.

Model	2K	4K	8K	16K	32K	64K	128K	256K
NoPE	99.2 → 98.8	75.8 → 88.7	48.9 → 81.1	47.5 → 46.7	35.2 → 35.5	13.2 → 16.6	5.1 → 9.2	0.6 → 0.9
Polar	98.9 → 98.6	92.4 → 97.8	87.5 → 97.2	73.1 → 77.6	67.8 → 83.5	54.3 → 69.0	41.8 → 63.6	34.4 → 47.1
RoPE	88.2 → 76.1	24.8 → 21.5	17.4 → 13.0	0.4 → 0.0	3.4 → 2.1	0.1 → 0.0	0.0 → 0.0	0.0 → 0.0

Table 23: Complete retrieval curve, untouched → capped, reported as teacher-forced target-token accuracy (%). Each entry averages passkey and NIAH, synthetic and FinePDFs haystacks, and three depths.

Model	2K	4K	8K	16K	32K	64K	128K	256K
NoPE	96.0 → 94.0	55.8 → 75.0	44.3 → 62.2	39.0 → 32.7	11.8 → 9.2	2.3 → 0.5	0.3 → 0.0	0.0 → 0.0
Polar	94.7 → 93.2	79.8 → 89.7	74.5 → 86.5	49.0 → 43.0	38.0 → 44.2	32.8 → 33.0	13.2 → 19.8	9.0 → 16.3
RoPE	59.8 → 45.0	2.2 → 0.8	0.3 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0

Table 24: Complete retrieval curve, untouched → capped, reported as exact five-token accuracy (%). Each entry averages passkey and NIAH, synthetic and FinePDFs haystacks, and three depths.

Model	Haystack	2K	4K	8K	16K	32K	64K	128K	256K
NoPE	Synthetic	99.1 → 98.7	98.7 → 97.7	97.7 → 95.8	94.9 → 90.8	70.3 → 69.7	26.5 → 33.2	10.2 → 18.4	1.2 → 1.9
	FinePDFs (Real)	99.3 → 98.8	52.9 → 79.8	0.0 → 66.4	0.0 → 2.7	0.0 → 1.3	0.0 → 0.1	0.0 → 0.0	0.0 → 0.0
Polar	Synthetic	99.1 → 98.5	98.4 → 97.5	97.1 → 96.1	97.0 → 96.0	94.9 → 93.9	92.3 → 92.5	80.3 → 85.5	68.4 → 77.5
	FinePDFs (Real)	98.6 → 98.7	86.3 → 98.1	77.9 → 98.3	49.1 → 59.1	40.7 → 73.1	16.3 → 45.5	3.2 → 41.6	0.3 → 16.7
RoPE	Synthetic	81.7 → 63.4	44.9 → 40.3	34.5 → 25.9	0.8 → 0.1	6.8 → 4.3	0.1 → 0.1	0.0 → 0.0	0.0 → 0.0
	FinePDFs (Real)	94.8 → 88.8	4.7 → 2.7	0.2 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0

Table 25: Teacher-forced retrieval token accuracy (%) by haystack type, untouched → capped. Polar gains 9.1 percentage points in synthetic and 16.3 points in FinePDFs retrieval at 256K, while NoPE FinePDFs retrieval remains 0.0% despite synthetic signal gains.

Model	Haystack	2K	4K	8K	16K	32K	64K	128K	256K
NoPE	Synthetic	95.3 → 93.7	93.7 → 88.7	88.7 → 79.7	78.0 → 65.3	23.7 → 18.3	4.7 → 1.0	0.7 → 0.0	0.0 → 0.0
	FinePDFs (Real)	96.7 → 94.3	18.0 → 61.3	0.0 → 44.7	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0
Polar	Synthetic	96.3 → 92.7	92.3 → 88.3	86.3 → 81.7	86.0 → 81.0	76.0 → 72.0	65.7 → 65.3	26.3 → 38.7	18.0 → 32.7
	FinePDFs (Real)	93.0 → 93.7	67.3 → 91.0	62.7 → 91.3	12.0 → 5.0	0.0 → 16.3	0.0 → 0.7	0.0 → 1.0	0.0 → 0.0
RoPE	Synthetic	42.3 → 27.0	4.3 → 1.7	0.7 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0
	FinePDFs (Real)	77.3 → 63.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0	0.0 → 0.0

Table 26: Exact five-token retrieval accuracy (%) by haystack type, untouched → capped. On real text, RoPE reaches 0.0% at 4K and NoPE at 8K. Polar scores 62.7% untouched and more than 91% capped at 8K, with nonzero capped accuracy through 128K.

These results support an isolated near-unit retention gate as a major mediator of NoPE and Polar extrapolation degradation. They motivate testing bounded retention gates during training, but the inference-only intervention does not validate that design.

The complete machine-readable artifacts are under `gamma_diagnostics/results/parameters/gamma_parameters.json` for the scan, `all_clamp_sweep_262k.json` for the paired pilot, and `re_evaluation/run-summary.json` plus per-suite

Model	0K	1K	2K	4K	8K	16K	32K	64K	128K	256K
NoPE	63 → 63	57 → 60	55 → 55	56 → 56	39 → 57	21 → 54	6 → 45	0 → 37	0 → 35	0 → 39
Polar	55 → 55	54 → 54	52 → 52	54 → 54	56 → 57	50 → 50	46 → 43	45 → 44	35 → 44	28 → 42
RoPE	59 → 61	67 → 68	64 → 66	43 → 50	28 → 33	19 → 21	15 → 17	16 → 15	14 → 13	14 → 11

Table 27: Complete BABILong curve after adaptation, untouched → capped (macro exact match, %). The cap raises NoPE to 35–57% across 8K–256K, compared with 39/21/6% untouched at 8/16/32K and zero from 64K onward, and raises Polar at 256K from 28% to 42%.

Model	2K	4K	8K	16K	32K	64K	128K	256K
NoPE	1.432 → 1.454	1.717 → 1.454	2.656 → 1.473	3.374 → 1.484	5.527 → 1.504	5.885 → 1.506	6.369 → 1.530	8.097 → 1.595
Polar	1.451 → 1.458	1.442 → 1.451	1.406 → 1.407	1.417 → 1.384	1.500 → 1.399	1.569 → 1.400	1.690 → 1.449	1.825 → 1.501
RoPE	1.439 → 1.440	1.642 → 1.656	1.783 → 1.783	1.935 → 1.932	2.094 → 2.082	2.252 → 2.220	2.413 → 2.378	2.512 → 2.460

Table 28: Complete fixed-target likelihood curve, untouched → capped, averaged over FinePDFs, PG-19, and Proof-Pile in mean bits per byte (lower is better).

Model	Dataset	2K	4K	8K	16K	32K	64K	128K	256K
NoPE	FinePDFs	0.913 → 0.924	0.868 → 0.885	1.366 → 0.900	1.473 → 0.920	5.317 → 0.969	5.519 → 1.010	6.637 → 1.060	11.666 → 1.180
	PG-19	1.197 → 1.211	1.461 → 1.216	2.191 → 1.207	2.473 → 1.200	4.259 → 1.205	4.799 → 1.207	4.927 → 1.214	4.902 → 1.224
	Proof-Pile	2.186 → 2.226	2.821 → 2.262	4.411 → 2.312	6.177 → 2.331	7.006 → 2.338	7.337 → 2.300	7.544 → 2.315	7.723 → 2.382
Polar	FinePDFs	0.946 → 0.950	0.913 → 0.913	0.929 → 0.925	0.938 → 0.901	1.032 → 0.924	1.113 → 0.967	1.252 → 1.011	1.419 → 1.052
	PG-19	1.226 → 1.229	1.224 → 1.227	1.221 → 1.215	1.229 → 1.212	1.259 → 1.222	1.287 → 1.243	1.331 → 1.277	1.374 → 1.319
	Proof-Pile	2.181 → 2.196	2.190 → 2.212	2.067 → 2.079	2.084 → 2.040	2.208 → 2.050	2.308 → 1.989	2.487 → 2.059	2.683 → 2.130
RoPE	FinePDFs	1.036 → 1.040	1.277 → 1.303	1.492 → 1.486	1.802 → 1.772	2.142 → 2.071	2.469 → 2.353	2.757 → 2.626	2.919 → 2.742
	PG-19	1.204 → 1.202	1.317 → 1.316	1.350 → 1.351	1.419 → 1.426	1.472 → 1.476	1.541 → 1.530	1.591 → 1.586	1.626 → 1.627
	Proof-Pile	2.077 → 2.078	2.331 → 2.350	2.507 → 2.513	2.582 → 2.596	2.668 → 2.700	2.746 → 2.776	2.892 → 2.922	2.991 → 3.011

Table 29: Dataset-level fixed-target long-document likelihood, untouched → capped, in bits per byte (lower is better). At 256K, the cap reduces NoPE BPB from 11.666 to 1.180 on FinePDFs, from 4.902 to 1.224 on PG-19, and from 7.723 to 2.382 on Proof-Pile.

logs for the full benchmark. The executable protocol and reporting caveats are in `gamma_diagnostics/README.md`.

K STRESS AND ENVIRONMENT DIAGNOSTICS

Figure 8 reports the sampled gains. For each block f_i , the diagnostic samples

$$G_i = \frac{\|f_i(x + \delta x) - f_i(x)\|_{\text{rms}}}{\|\delta x\|_{\text{rms}}}, \quad \|\delta x\|_{\text{rms}} = 0.02\|x\|_{\text{rms}}. \quad (13)$$

The maximum sampled gain is a lower bound on the local operator norm along sampled directions. It is not a tight bound on the block Lipschitz constant. In the observed NoPE collapse, changes first concentrate near blocks 6–7 and later propagate through the residual stream. Polar’s sampled gains and activation moments remain comparatively flat through 256K.

L CHECKPOINT-VARIABILITY DIAGNOSTIC JOURNEY

The variability finding emerged sequentially. Stage I used L4 GPUs, followed by the first 10B NoPE control on an L4. The difference became visible after primary training moved to L40S. Short-context convergence remained close (2.8126 vs 2.8160 nats; approximately 0.12% difference), while 256K NLL differed (4.07 vs 10.77 nats). Table 30 records the seven-step diagnostic sequence and the retention-horizon intervention.

This sequence localizes the observed checkpoint failure to one unusually long zero-input retention horizon, and the inference cap reverses much of the measured degradation. It does not establish what training event produced the outlier or whether hardware-level differences contributed to it. All five headline Stage II models were trained in the same L40S environment, providing an internally controlled comparison for those runs.

FEA Structural Stress Probing: Local Block Secant Gains G Across Length (2K -> 256K)

Polar (ATMA): Flat Gains Across Length									RoPE: Representational Attenuation									NoPE (L40S mbs16): Block 7 Yield Locus								
	2K	4K	8K	16K	32K	64K	128K	256K		2K	4K	8K	16K	32K	64K	128K	256K		2K	4K	8K	16K	32K	64K	128K	256K
B0	0.86	0.86	0.86	0.86	0.86	0.86	0.86	0.86	B0	0.94	0.94	0.94	0.94	0.94	0.94	0.94	0.94	B0	0.91	0.91	0.91	0.91	0.91	0.91	0.91	0.91
B1	0.78	0.78	0.78	0.78	0.78	0.78	0.78	0.78	B1	0.92	0.92	0.92	0.92	0.92	0.92	0.92	0.92	B1	0.88	0.88	0.88	0.88	0.88	0.88	0.88	0.88
B2	0.96	0.95	0.95	0.95	0.96	0.96	0.96	0.96	B2	1.01	1.00	0.99	0.98	0.98	0.98	0.98	0.98	B2	1.02	1.02	1.02	1.02	1.02	1.03	1.03	1.03
B3	0.90	0.89	0.89	0.89	0.88	0.88	0.88	0.87	B3	1.02	1.02	1.02	1.02	1.02	1.02	1.02	1.02	B3	0.98	0.98	0.98	0.97	0.97	0.97	0.97	0.97
B4	0.87	0.87	0.86	0.86	0.86	0.85	0.85	0.85	B4	1.05	1.07	1.08	1.07	1.06	1.04	1.03	1.03	B4	0.96	0.96	0.96	0.96	0.96	0.96	0.96	0.95
B5	0.86	0.85	0.85	0.84	0.84	0.84	0.84	0.84	B5	0.94	0.93	0.93	0.93	0.93	0.93	0.93	0.92	B5	0.94	0.94	0.94	0.94	0.95	0.96	0.96	0.96
B6	1.02	1.02	1.03	1.05	1.07	1.09	1.10	1.11	B6	1.09	1.06	1.04	1.02	1.01	1.01	1.00	1.00	B6	1.13	1.11	1.09	1.09	1.08	1.09	1.09	1.10
B7	0.98	0.97	0.97	0.97	0.97	0.97	0.97	0.97	B7	1.01	1.01	1.00	1.00	1.00	1.00	0.99	0.99	B7	1.02	1.02	1.02	1.03	1.09	1.20	1.31	1.38
B8	0.96	0.96	0.95	0.95	0.95	0.95	0.95	0.94	B8	1.01	1.00	1.00	1.00	0.99	0.99	0.99	0.99	B8	1.01	1.01	1.01	1.00	1.00	1.00	1.00	0.99
B9	0.93	0.92	0.92	0.92	0.92	0.92	0.92	0.91	B9	0.98	0.98	0.98	0.97	0.97	0.97	0.97	0.96	B9	0.99	0.99	0.99	0.98	0.98	0.98	0.98	0.99
B10	1.10	1.10	1.10	1.10	1.10	1.10	1.10	1.09	B10	1.09	1.08	1.07	1.06	1.06	1.05	1.04	1.04	B10	1.16	1.14	1.12	1.12	1.12	1.11	1.11	1.11
B11	1.02	1.02	1.02	1.02	1.02	1.01	1.01	1.01	B11	1.02	1.02	1.02	1.01	1.01	1.01	1.00	1.00	B11	1.04	1.03	1.03	1.02	1.01	1.01	1.00	1.00
B12	1.02	1.01	1.01	1.01	1.01	1.00	1.00	1.00	B12	1.02	1.02	1.01	1.01	1.01	1.00	1.00	1.00	B12	1.04	1.04	1.03	1.02	1.01	1.01	1.01	1.00
B13	1.01	1.01	1.00	1.00	1.00	1.00	1.00	1.00	B13	1.02	1.01	1.01	1.00	1.00	1.00	1.00	1.00	B13	1.04	1.04	1.06	1.08	1.09	1.07	1.06	1.06
B14	1.43	1.44	1.45	1.44	1.41	1.39	1.38	1.37	B14	1.32	1.31	1.29	1.28	1.27	1.26	1.25	1.25	B14	1.39	1.37	1.35	1.35	1.30	1.26	1.24	1.23
B15	1.23	1.23	1.22	1.21	1.20	1.19	1.18	1.17	B15	1.21	1.20	1.19	1.19	1.18	1.17	1.17	1.16	B15	1.18	1.18	1.17	1.14	1.12	1.11	1.09	1.08

Figure 8: Block-wise sampled secant gains for the three matched attention checkpoints over 64 nested FinePDFs documents and 64 perturbation directions per document. Raven Native and Atma-Raven-Titans were not run under this checkpoint-specific probe. The diagnostic localizes changes but does not estimate a global Lipschitz constant.

Step	Observation or intervention	Diagnostic implication
1. L4 baseline	Stage I and first 10B NoPE control on L4 (mbs4); 256K NLL was 4.07.	Established baseline extrapolation, but lacked replicated device-class controls.
2. L40S transition	Promoted NoPE run on L40S (mbs16); 2K validation matched (2.8126 vs 2.8160), but 256K NLL rose to 10.77.	Shown that the short-context pretraining curves did not predict long-context behavior.
3. Evaluation replay	L4 and L40S checkpoints re-evaluated under PyTorch 2.12 and 2.13 with identical metrics.	Excluded evaluation software runtime mismatch as the source.
4. Microbatch control	NoPE retrained on L40S with mbs4; 2K validation ended at 2.8128 and 256K NLL at 9.90.	Shown that microbatch size four did not reliably prevent collapse; initialization remained uncontrolled.
5. Deterministic fixtures	Synthetic 2K stream fixtures across device classes showed zero gradient or loss divergence.	Confirmed short-context kernel execution consistency; left training trajectory unconstrained.
6. Retention parameter inspection	Evaluated learned zero-input operating point $\gamma_0 = \sigma(b_\gamma)$ across all layers and heads.	Identified an outlier: L40S NoPE developed an isolated Block 2 head with $b_\gamma = +13.33$ ($H_{1/2} = 21.01\text{M}$ tokens), whereas L4 NoPE had max $H_{1/2} = 185.3$ tokens.
7. Single-head causal intervention	Installed a reversible inference cap ($H_{\max} = 256$ tok) on the single outlier head and re-evaluated 2K–256K.	Tested the mechanism: Capping the outlier head reduced 256K NLL from 12.00 to 1.14 nats and raised 256K BABILong to 39%, supporting excessive retention as a mediator of the collapse.

Table 30: Diagnostic sequence from the initial L4/L40S divergence to the retention-horizon intervention. Steps 1–5 tested evaluation artifacts and showed that reducing microbatch size did not reliably prevent the failure. Steps 6–7 identified the isolated retention-gate outlier and tested its effect.